



**World's Leading Digital
Talent Transformation
Company**



World's Leading Digital Talent Transformation Company



130,000+ Professionals
Trained world-wide



Delivered Training in over
45 Countries



Over 7000 Industry
Veterans as Instructors



450+ Customers

Authorized Training Partner For



Red Hat



AUTOMATION
ANYWHERE



Silver
**Microsoft
Partner**



WORLD
HRD
CONGRESS

Cognixia is awarded as
Training Company of the Year, 2018



Asian Training &
Leadership Award, Dubai



ISO 9001:2015 Certified
Quality Management System



ISO/IEC 27001:2013 Certified
Information Security Management System

Centres of Excellence



**Machine
Learning & AI**



Microservices



**Robotic Process
Automation**



**Internet of
Things**



**Big Data
Analytics**



DevOps



**Cloud
Computing**



**Cyber
Security**



**BI
Technology**



**Professional
Development**

Upcoming Master Classes

25th April



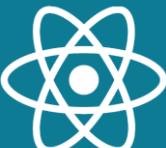
Machine Learning
in Production

9th May



Azure IoT

23rd May



React.js

30th May



Amazon Web
Services

Live Instructor Led Virtual Classes

Data Science

Machine Learning

AWS

Azure

Python

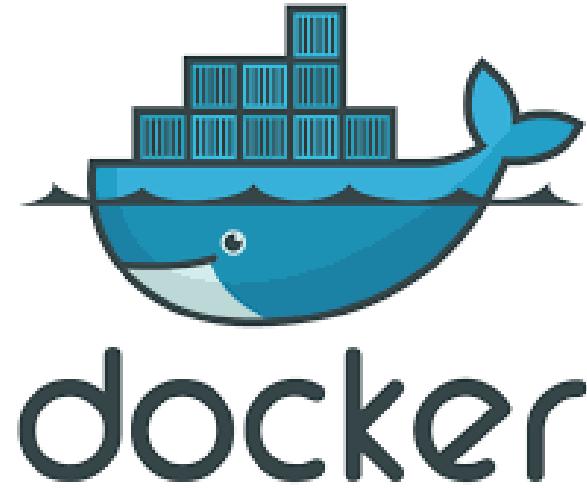
Spring Boot &
Microservices

ITIL v4 Foundation



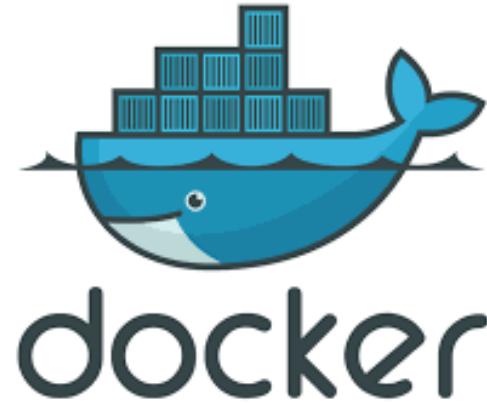
Docker & Kubernetes Workshop





Docker Overview

What is Docker?

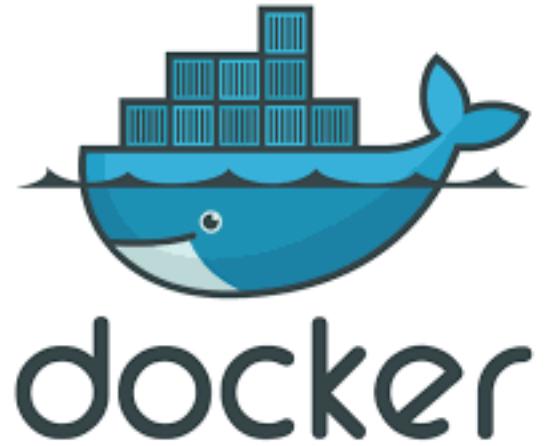


Docker is an open-source project that automates the deployment of applications inside software containers, by providing an additional layer of abstraction and automation of operating system-level virtualization on Linux.

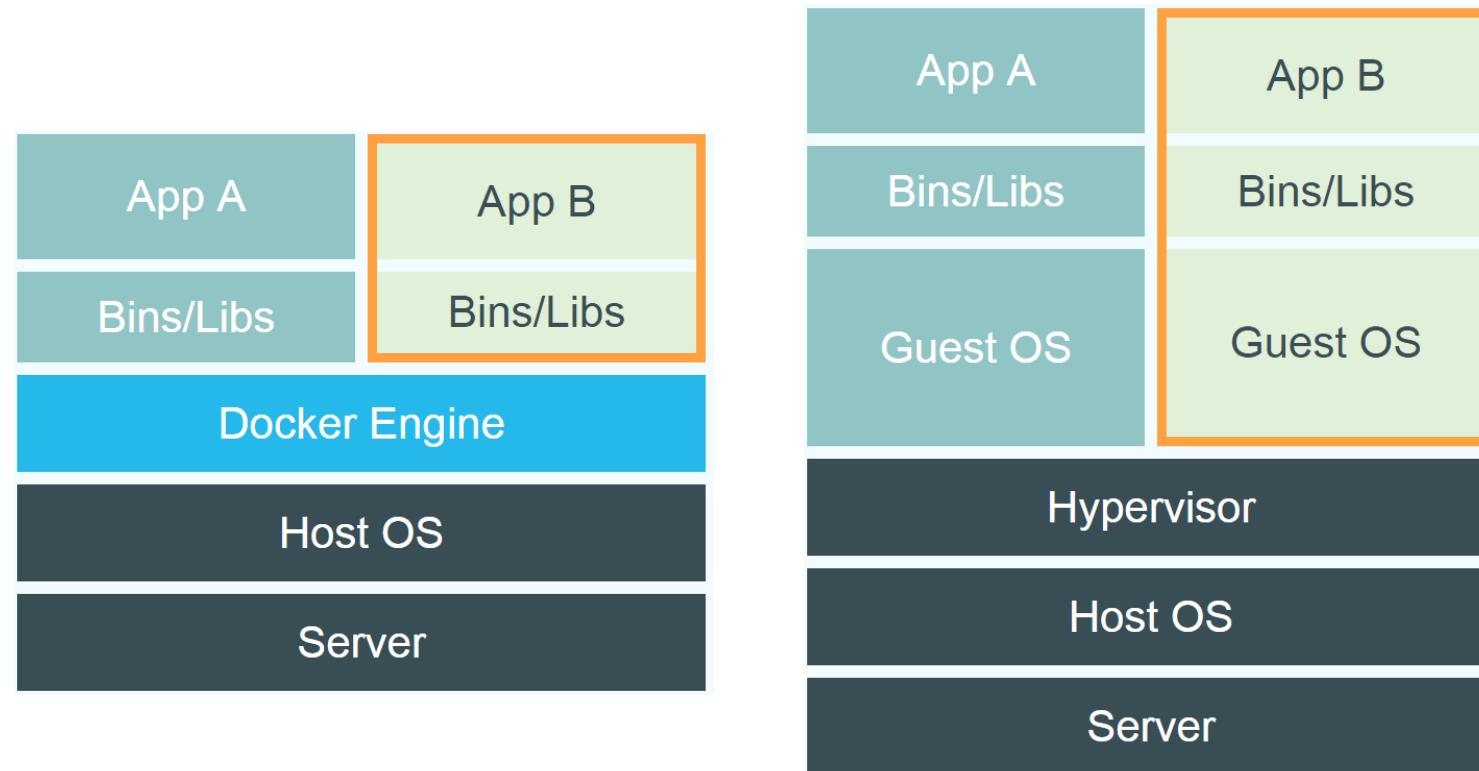
[Source: en.wikipedia.org]

Docker: Name

- Provide a uniformed wrapper around a software package: «*Build, Ship and Run Any App, Anywhere*»
- Similar to shipping containers: The container is always the same, regardless of the contents and thus fits on all trucks, cranes, ships, ...

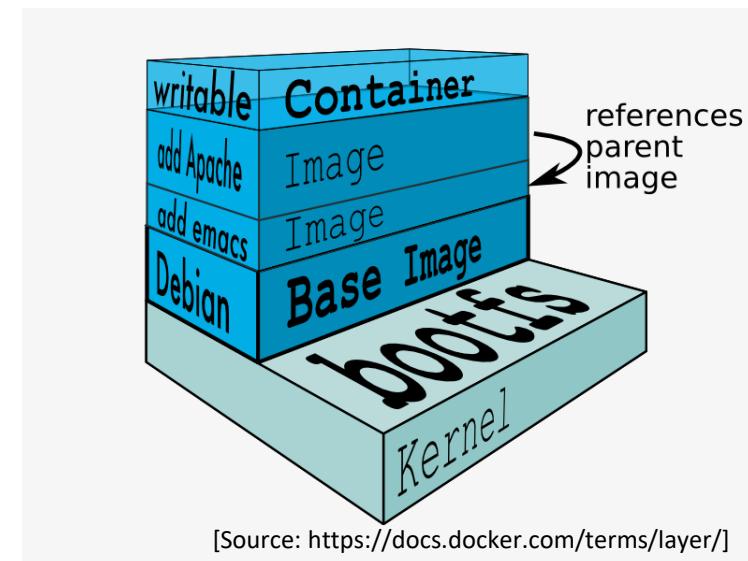


Docker vs. Virtual Machine



Source: <https://www.docker.com/whatisdocker/>

- libvirt: Platform Virtualization
- LXC (LinuX Containers): Multiple isolated Linux systems (containers) on a single host
- Layered File System

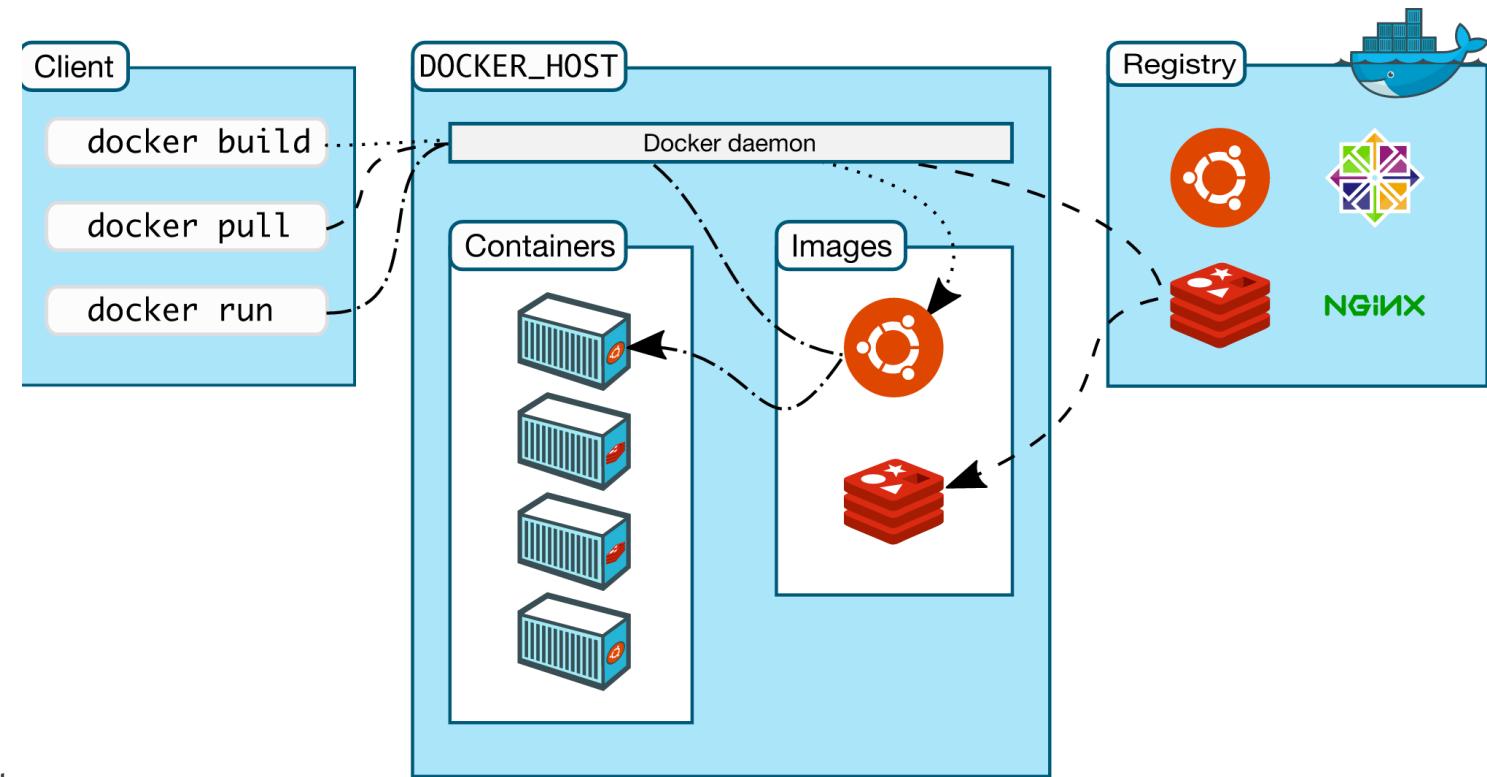


Containers

- Containers are operating system level virtualization
 - Allows for multiple isolated user space instances called containers
 - They share a single kernel
 - Can be added or removed at any time
- Containers consist of a self contained Linux file system
 - Can be from any Linux distribution which is compatible with the host kernel
 - Usually contain a single application such as a server
- Operating system level virtualization is lightweight
 - Is often used in Cloud Computing

Docker Architecture

- Docker uses a client-server architecture
- Client
 - Is the primary user interface which communicates using a REST API
 - Over HTTP
 - Over local Unix socket
- Server
 - Is the Docker daemon
 - Responsible for building, running, and distributing containers
- Registry
 - Responsible for the storage, management, and delivery of Docker Images
 - Docker Hub
 - Private
 - Other vendors



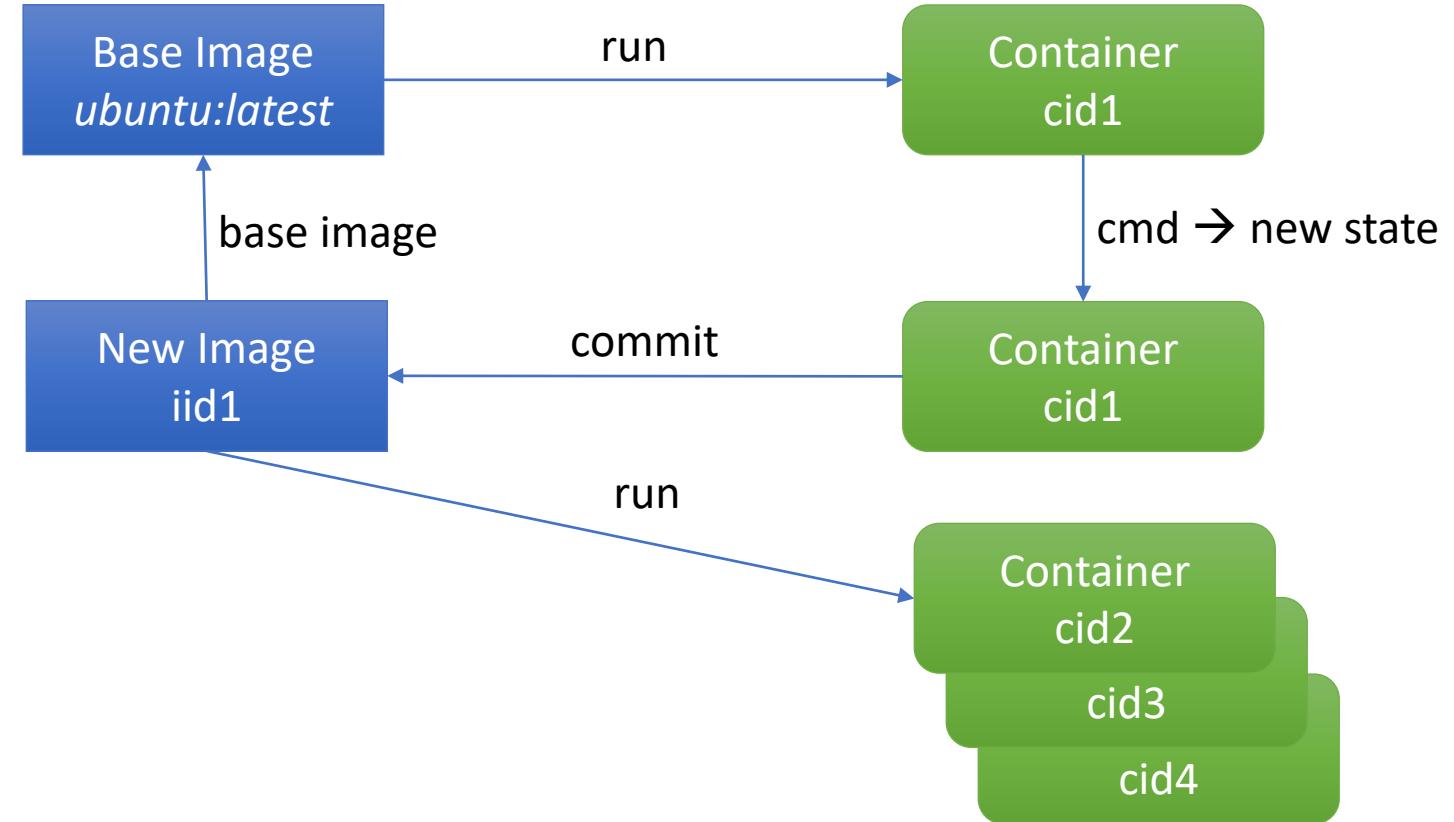
Terminology - Image

- Persisted snapshot that can be run
 - *images*: List all local images
 - *run*: Create a container from an image and execute a command in it
 - *tag*: Tag an image
 - *pull*: Download image from repository
 - *rmi*: Delete a local image
 - This will also remove intermediate images if no longer used

Terminology - Container

- Runnable instance of an image
 - *ps*: List all running containers
 - *ps -a*: List all containers (incl. stopped)
 - *top*: Display processes of a container
 - *start*: Start a stopped container
 - *stop*: Stop a running container
 - *pause*: Pause all processes within a container
 - *rm*: Delete a container
 - *commit*: Create an image from a container

Image vs. Container



- Create images automatically using a build script: «Dockerfile»
- Can be versioned in a version control system like Git or SVN, along with all dependencies
- Docker Hub can automatically build images based on dockerfiles on Github

Dockerfile Example

- Dockerfile:
 - FROM busybox:latest
MAINTAINER Your Name (email@domain.com)
CMD ["date"]
 - docker build [DockerFileDir]
 - docker inspect [imageId]

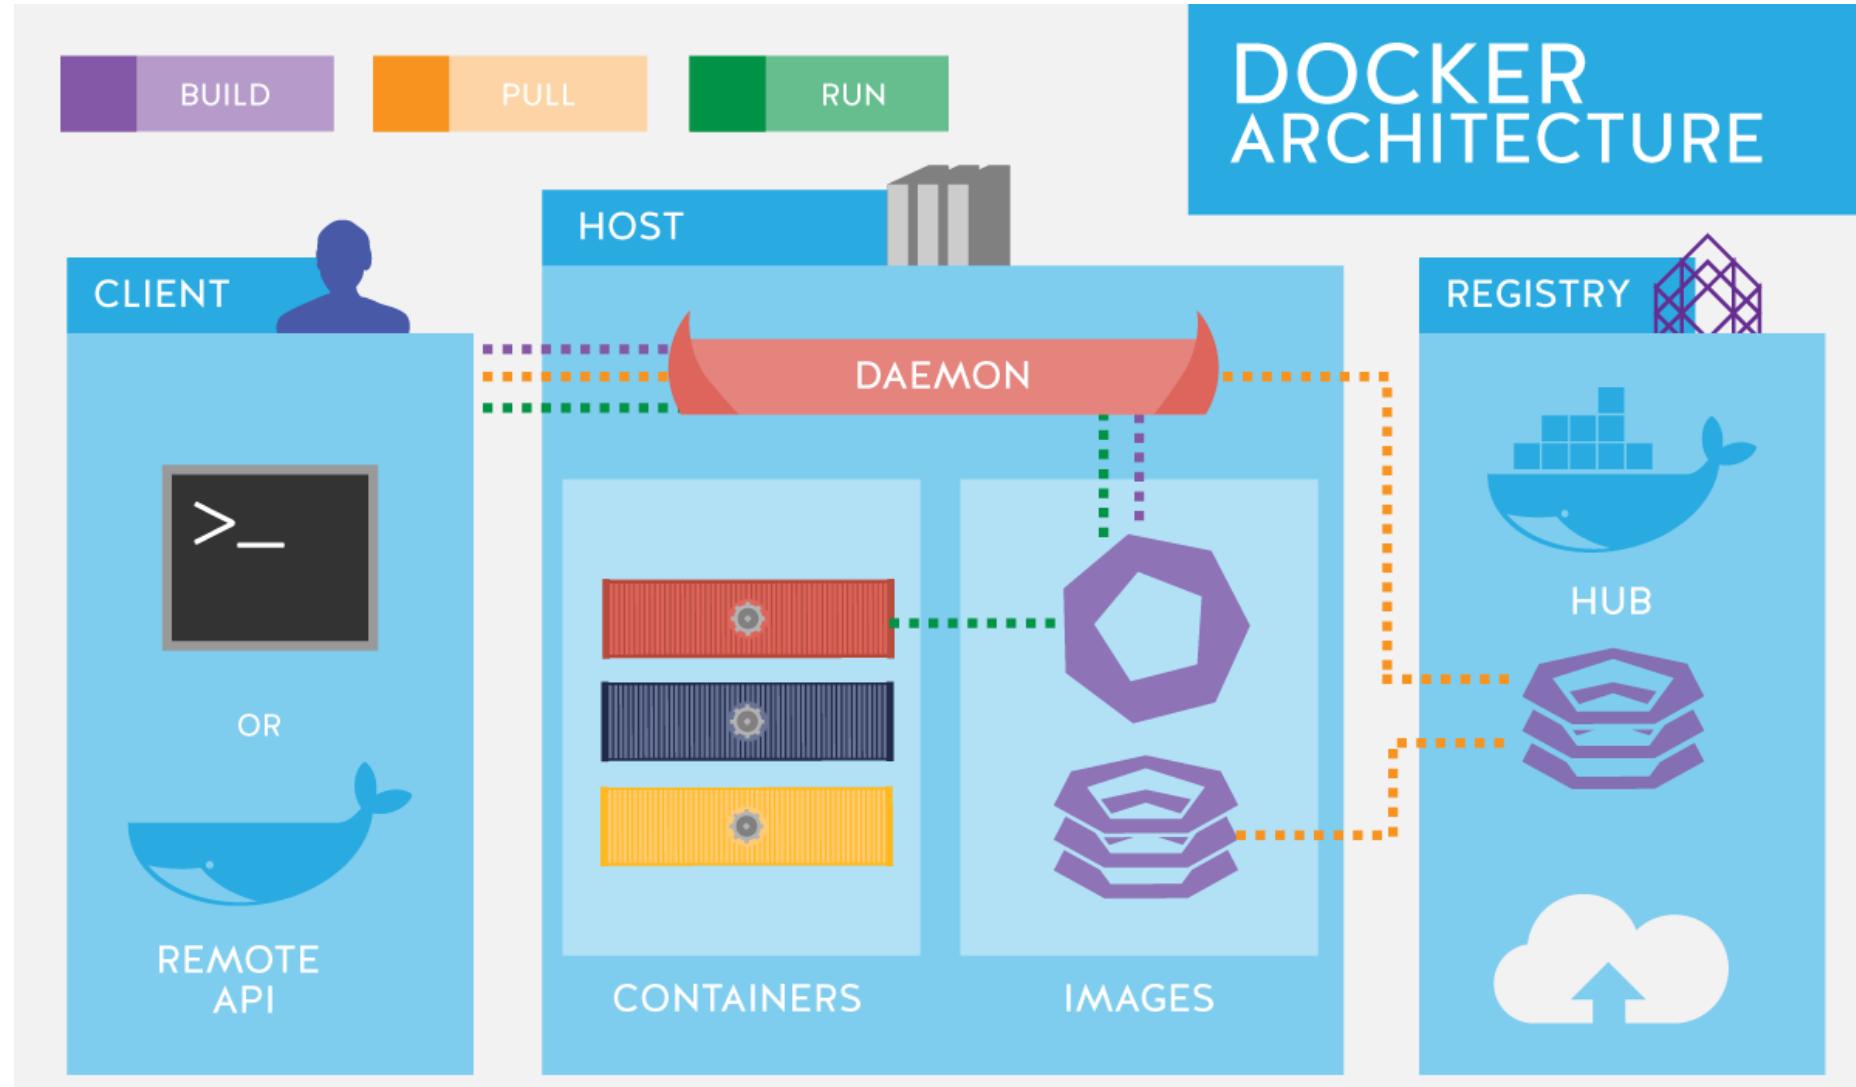
Mount Volumes

- docker run -ti **-v /hostLog:/log** ubuntu
- Run second container: Volume can be shared
 - docker run -ti --volumes-from firstContainerName ubuntu

Publish Port

- docker run -t -p 8080:80 ubuntu nc -l 80
 - Map container port 80 to host port 8080
 - Check on host: nc localhost 8080

Docker Architecture

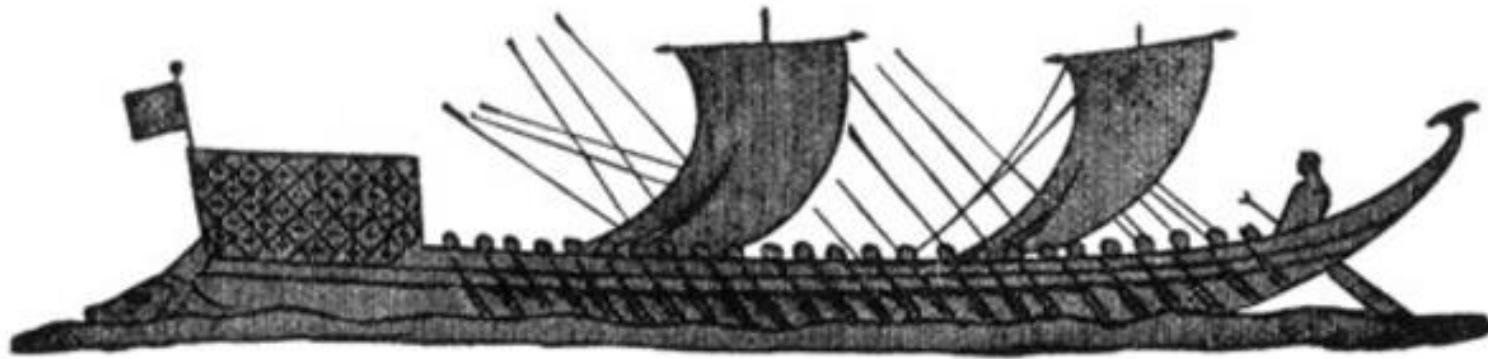




Kubernetes Overview

What Does "Kubernetes" Mean?

- Greek for “pilot” or “Helmsman of a ship”



What is Kubernetes?

- Project that was spun out of Google as an open source container orchestration platform.
- Built from the lessons learned in the experiences of developing and running Google's Borg and Omega.
- Designed from the ground-up as a loosely coupled collection of components centered around deploying, maintaining and scaling workloads.

What Does Kubernetes do?

- Known as the linux kernel of distributed systems.
- Abstracts away the underlying hardware of the nodes and provides a uniform interface for workloads to be both deployed and consume the shared pool of resources.
- Works as an engine for resolving state by converging actual and the desired state of the system.

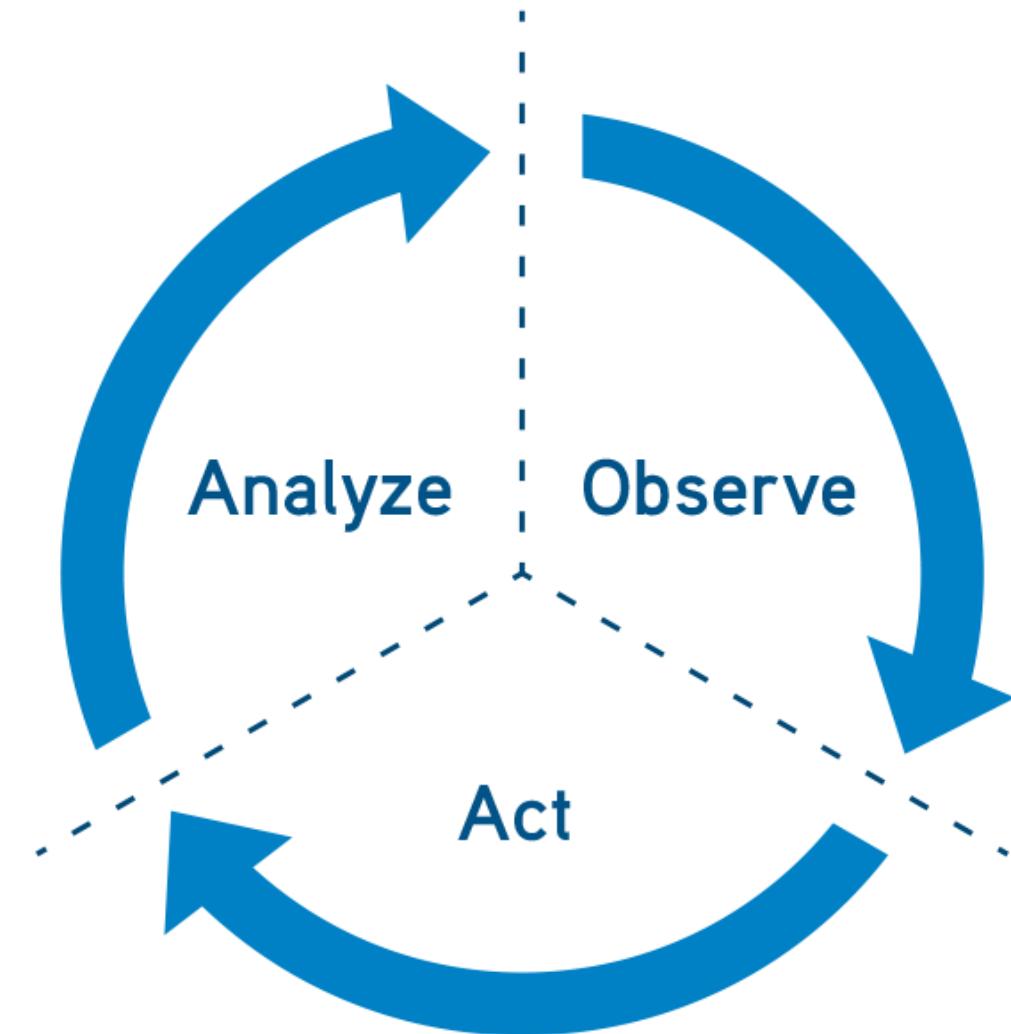
Decouples Infrastructure and Scaling

- All services within Kubernetes are natively Load Balanced.
- Can scale up and down dynamically.
- Used both to enable self-healing and seamless upgrading or rollback of applications.

Self Healing

- Kubernetes will **ALWAYS** try and steer the cluster to its desired state.
 - **Me:** “I want 3 healthy instances of redis to always be running.”
 - **Kubernetes:** “Okay, I’ll ensure there are always 3 instances up and running.”
 - **Kubernetes:** “Oh look, one has died. I’m going to attempt to spin up a new one.”

Kubernetes reconciliation loop



What can Kubernetes REALLY do?

- Autoscale Workloads
- Blue/Green Deployments
- Fire off jobs and scheduled cronjobs
- Manage Stateless and Stateful Applications
- Provide native methods of service discovery
- Easily integrate and support 3rd party apps

Most Importantly...

Use the **SAME API**
across bare metal and
EVERY cloud provider!!!

Who "Manages" Kubernetes?



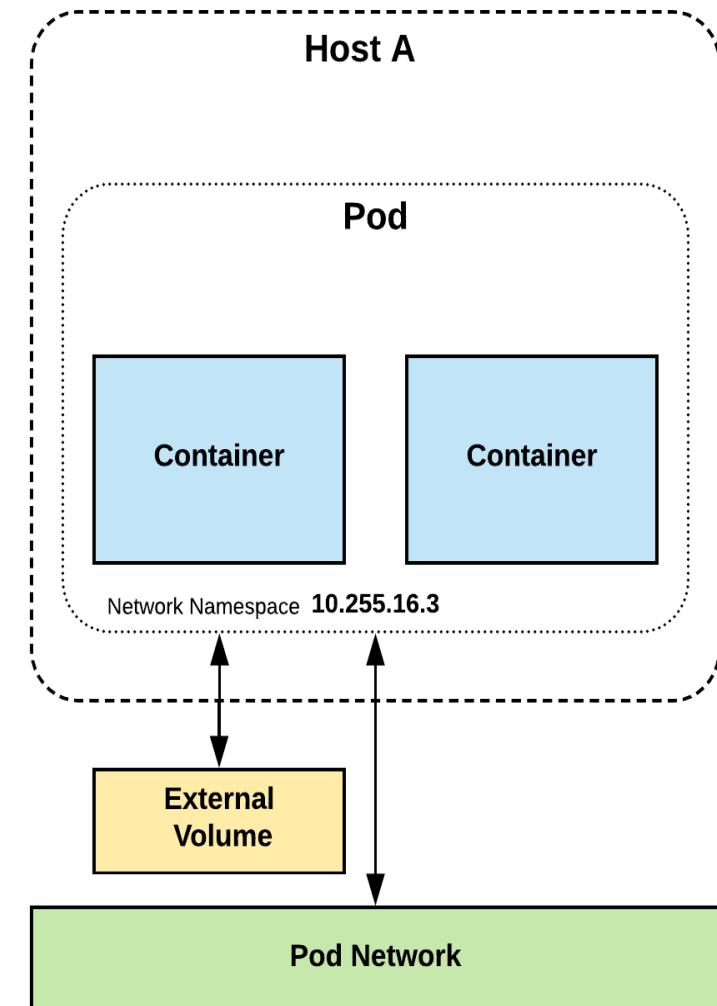
CLOUD NATIVE COMPUTING FOUNDATION

- The CNCF is a child entity of the Linux Foundation and operates as a vendor neutral governance group.

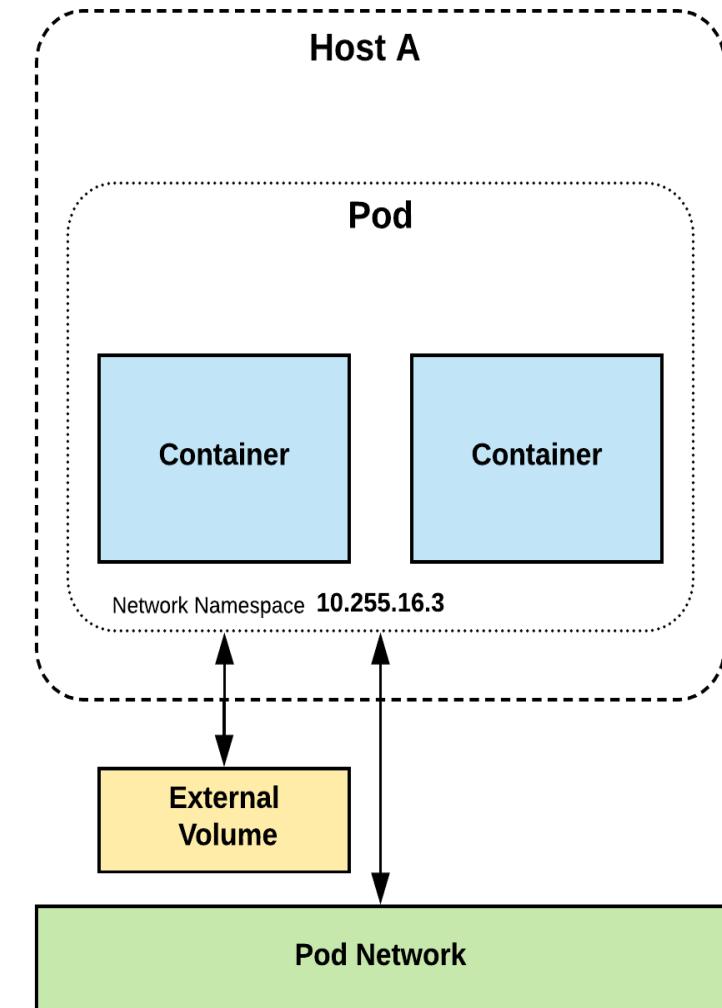
A Couple Key Concepts...

Pods

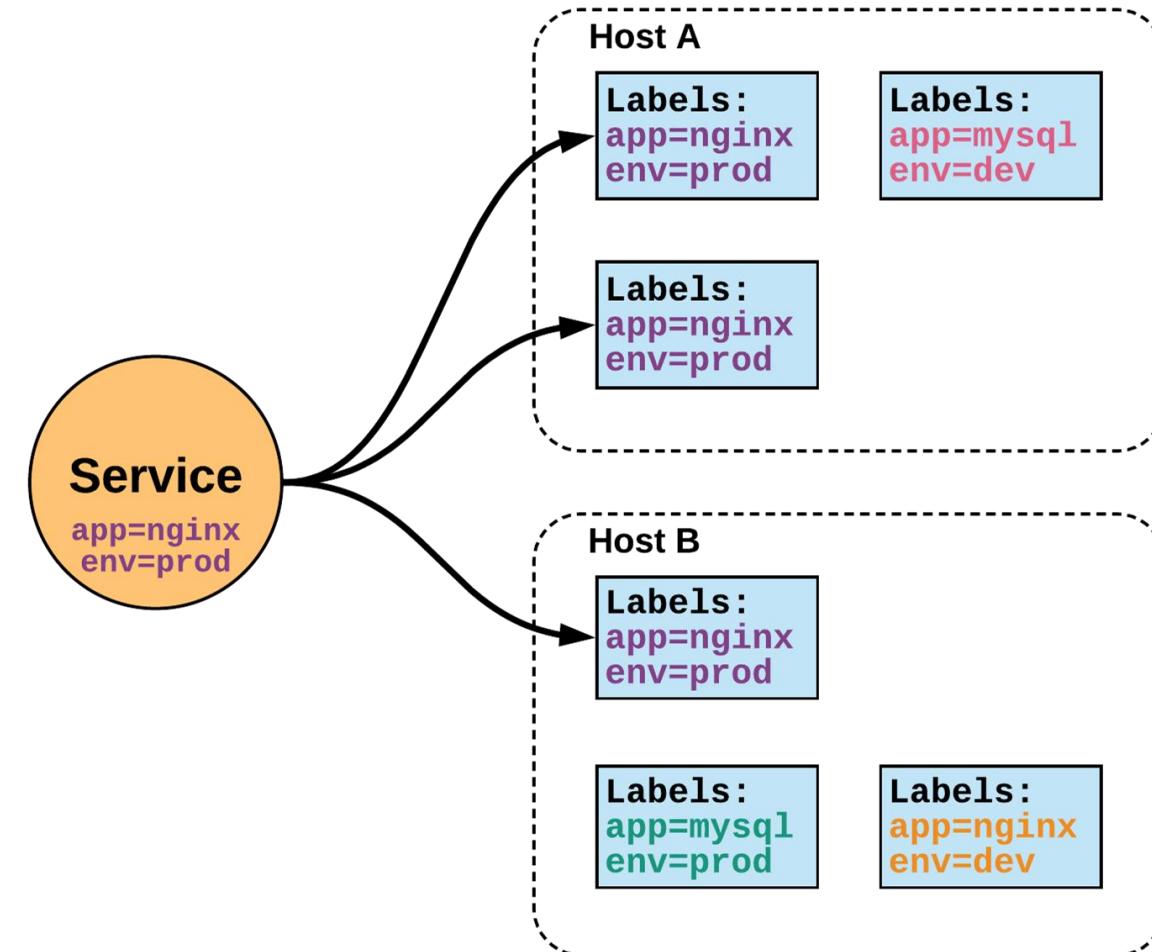
- **Atomic unit** or smallest “*unit of work*” of Kubernetes.
- Pods are **one or MORE containers** that share volumes, a network namespace, and are a part of a **single context**.



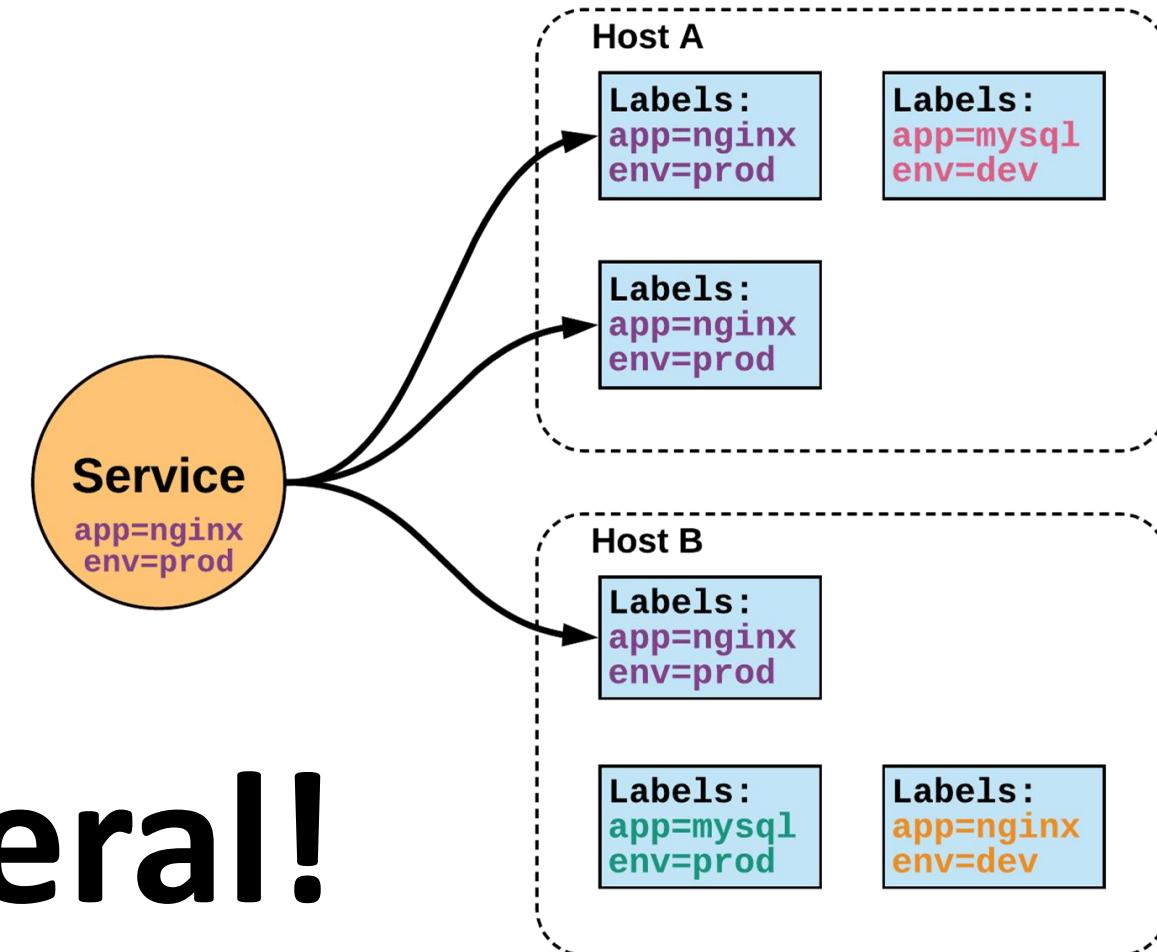
They are also Ephemeral!



- **Unified method of accessing the exposed workloads of Pods.**
- **Durable resource**
 - static cluster IP
 - static namespaced DNS name



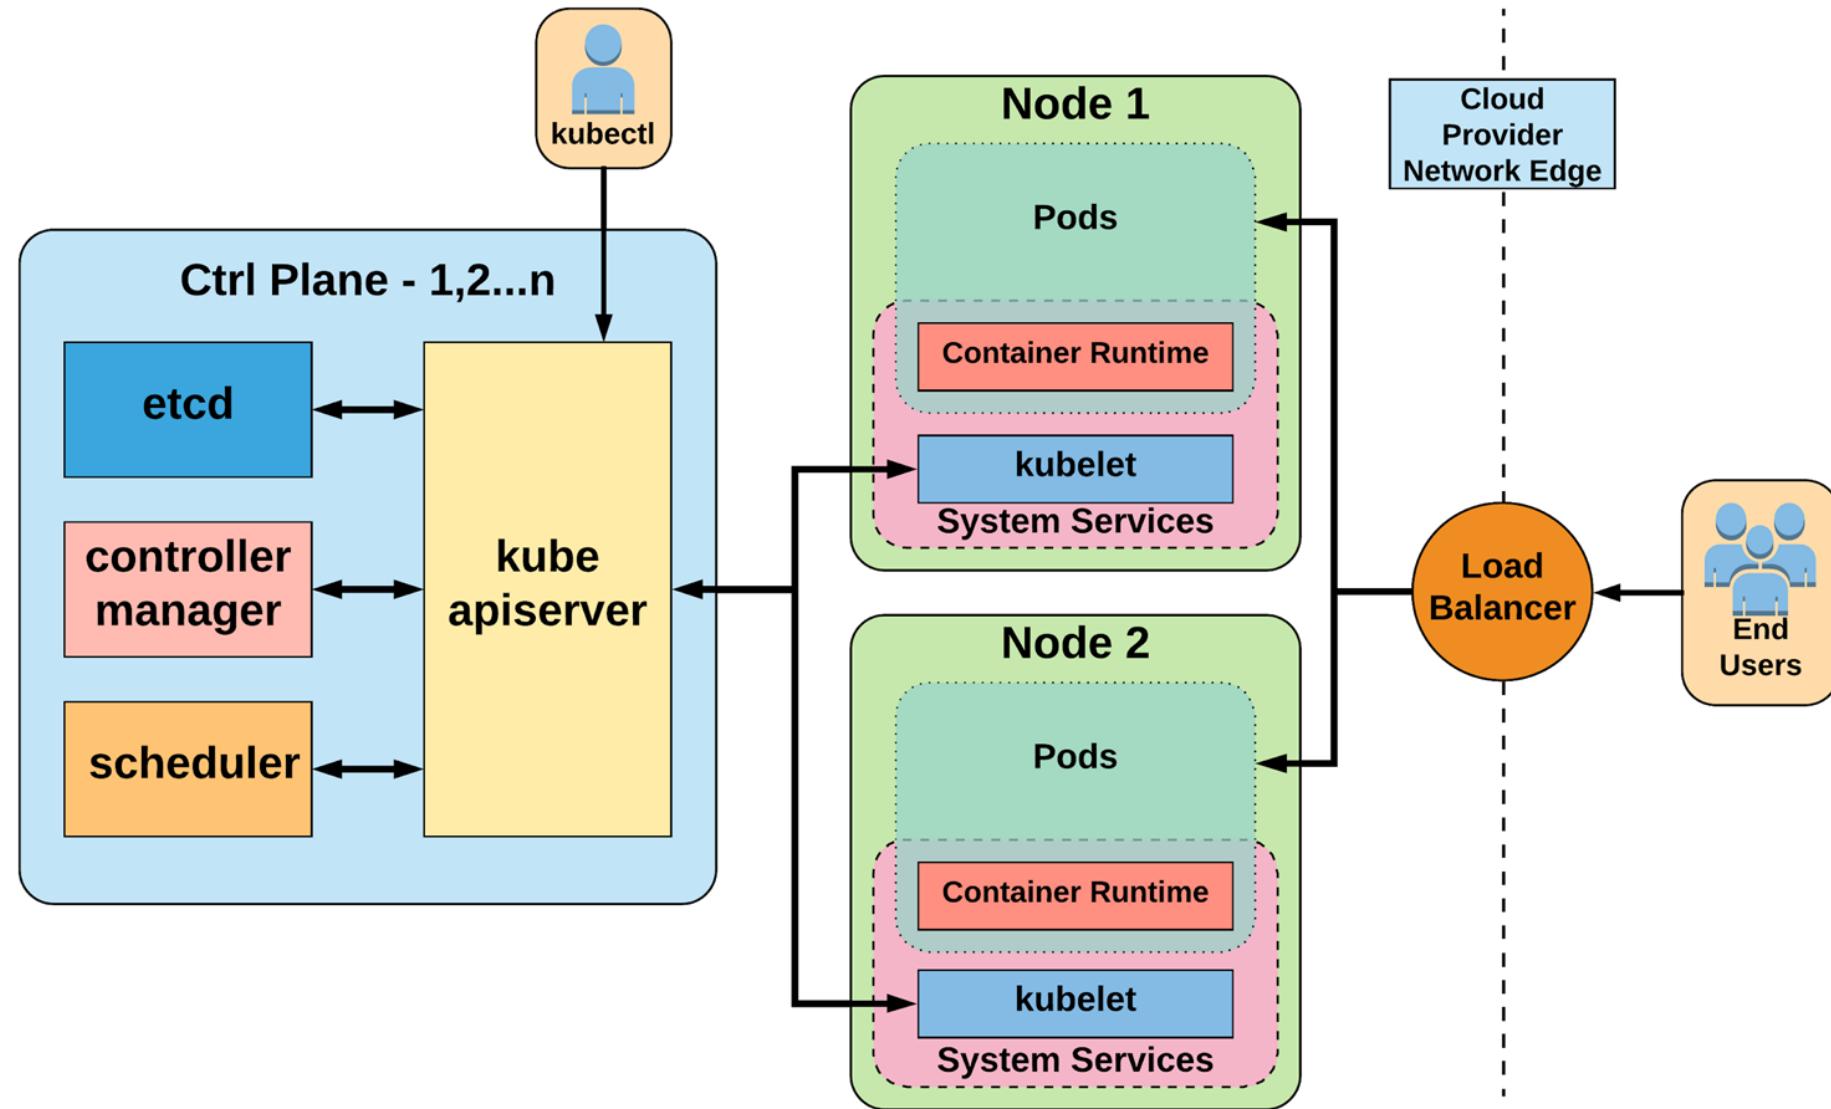
- **Unified method of accessing the exposed workloads of Pods.**
- **Durable resource**
 - static cluster IP
 - static namespaced DNS name



NOT Ephemerall!

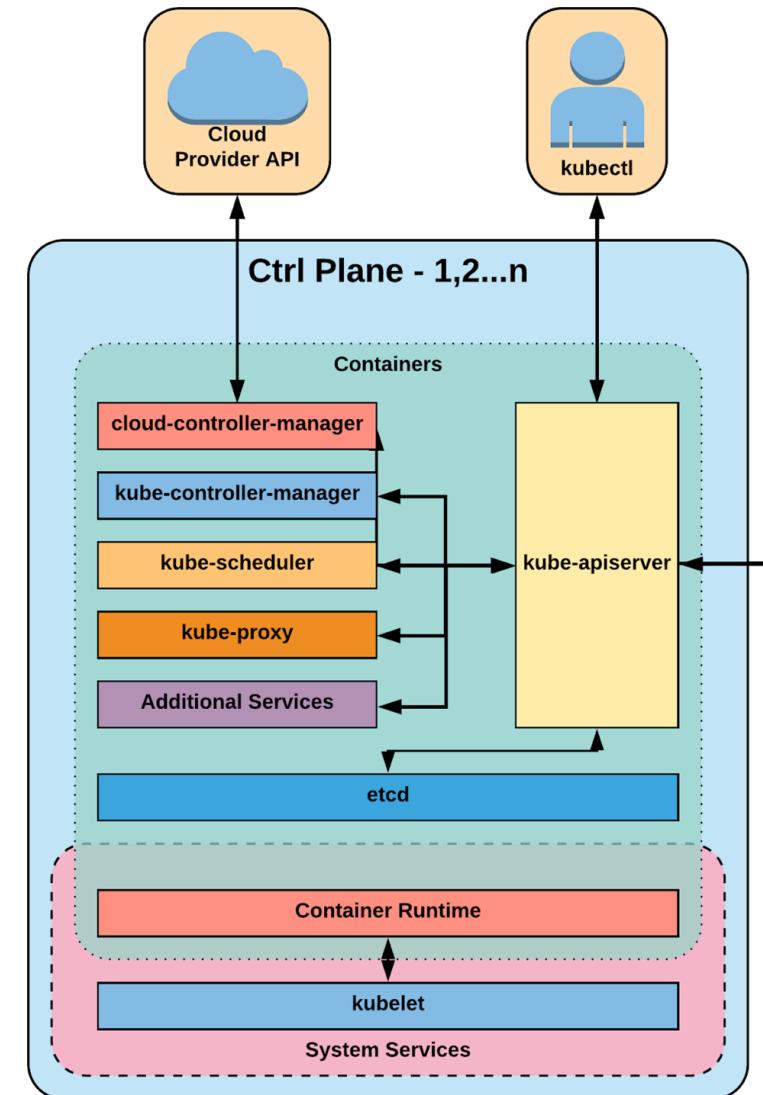
Architecture Overview

Control Plane Components – Architecture Overview



Control Plane Components

- **kube-apiserver**
- **etcd**
- **kube-controller-manager**
- **kube-scheduler**



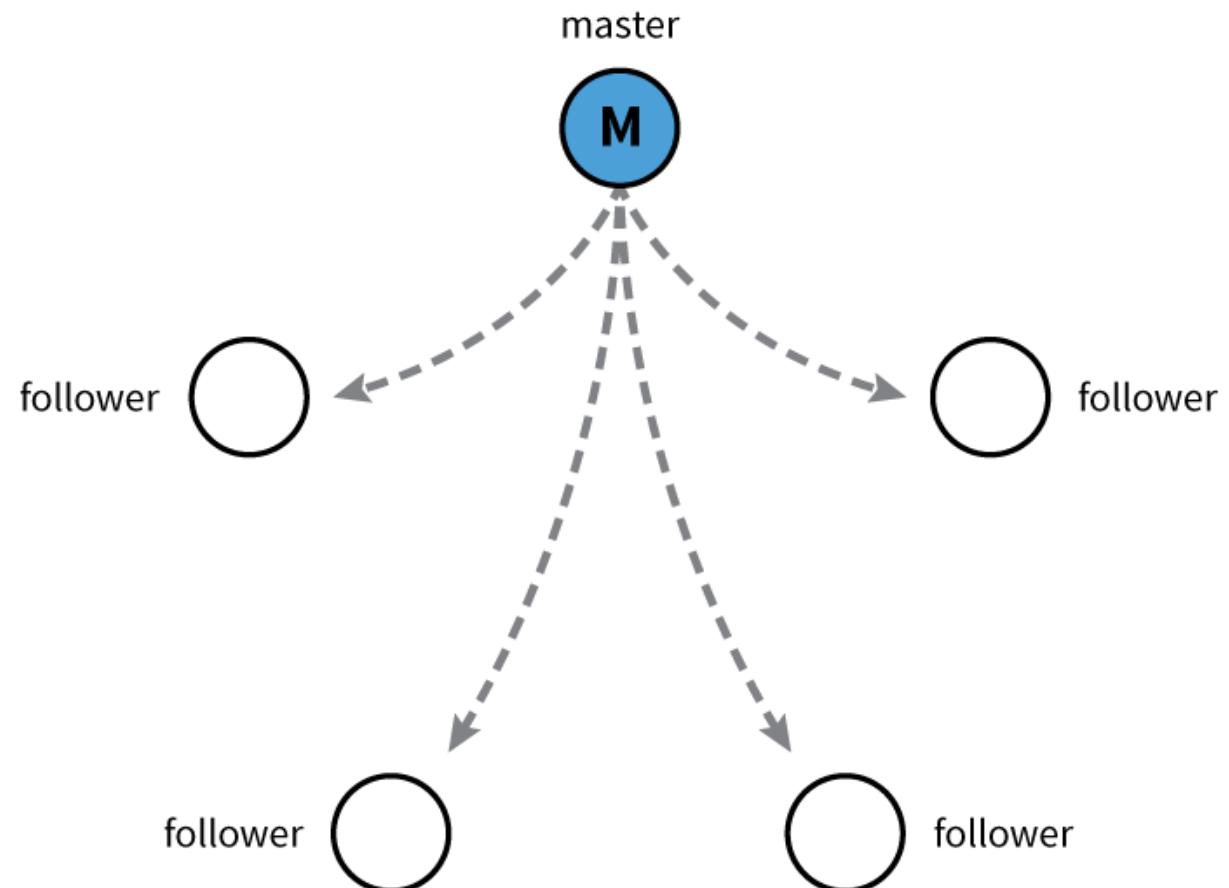
kube-apiserver

- Provides a forward facing **REST** interface into the kubernetes control plane and datastore.
- All clients and other applications interact with kubernetes **strictly** through the API Server.
- Acts as the gatekeeper to the cluster by handling authentication and authorization, request validation, mutation, and admission control in addition to being the front-end to the backing datastore.

- **etcd** acts as the cluster datastore.
- Purpose in relation to Kubernetes is to provide a strong, consistent and highly available key-value store for persisting cluster state.
- Stores objects and config information.



- Uses “Raft Consensus” among a quorum of systems to create a fault-tolerant consistent “view” of the cluster.
- <https://raft.github.io/>



kube-controller-manager

- Serves as the primary daemon that manages all core component control loops.
- Monitors the cluster state via the apiserver and **steers the cluster towards the desired state.**

List of core controllers:

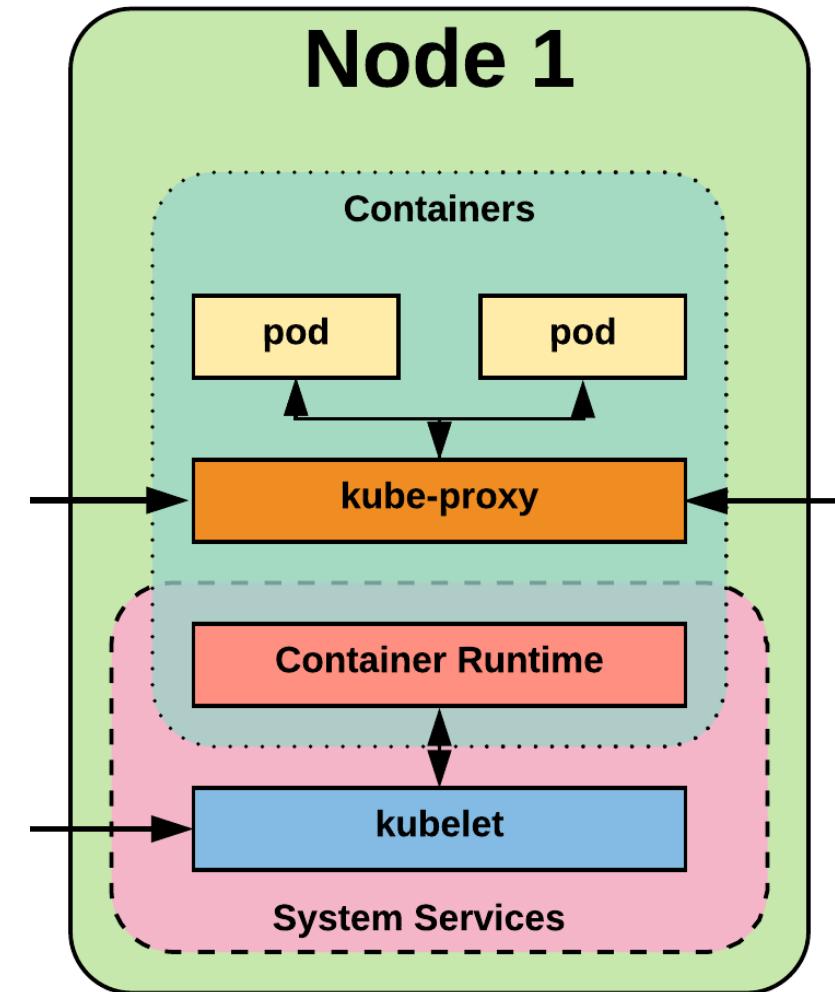
<https://github.com/kubernetes/kubernetes/blob/master/cmd/kube-controller-manager/app/controllermanager.go#L344>

kube-scheduler

- Verbose policy-rich engine that evaluates workload requirements and attempts to place it on a matching resource.
- Default scheduler uses bin packing.
- Workload Requirements can include: general hardware requirements, affinity/anti-affinity, labels, and other various custom resource requirements.

Node Components – Architecture Overview

- **kubelet**
- **kube-proxy**
- **Container Runtime Engine**



kubelet

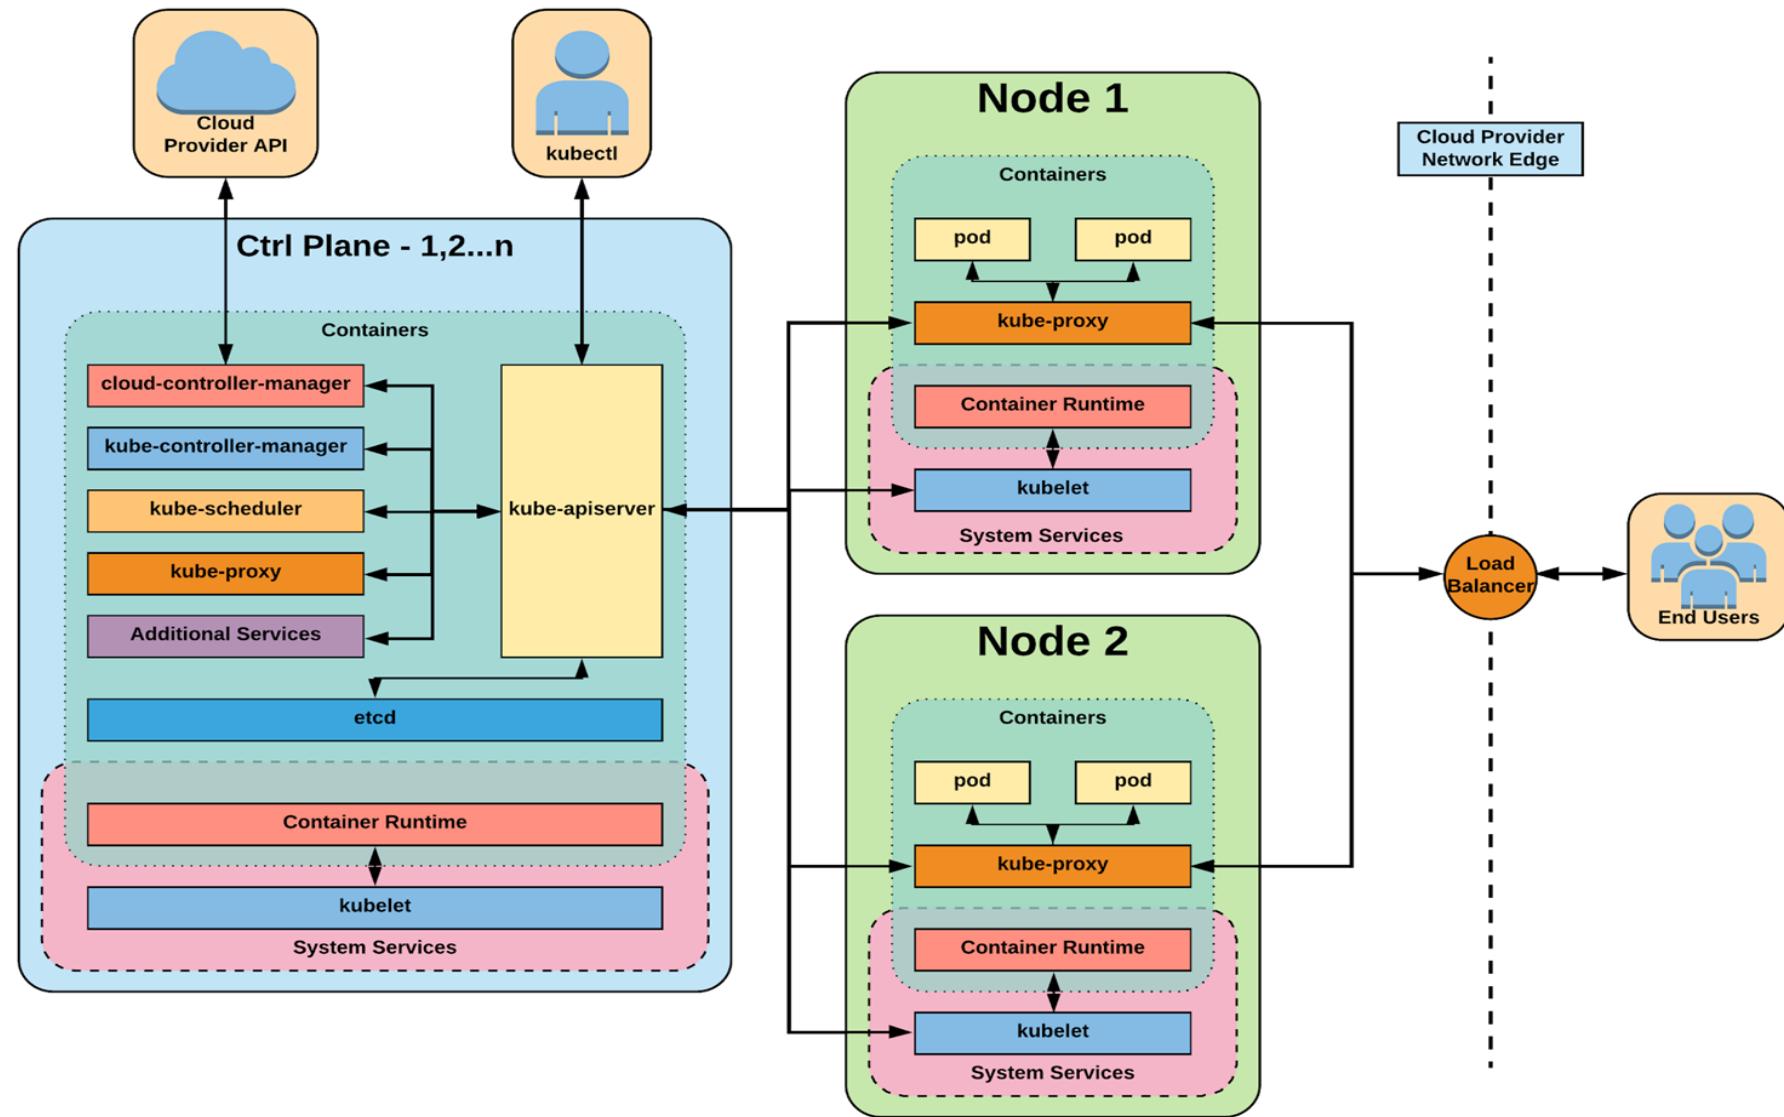
- Acts as the node agent responsible for managing the lifecycle of every pod on its host.
- Kubelet understands YAML container manifests that it can read from several sources:
 - file path
 - HTTP Endpoint
 - etcd watch acting on any changes
 - HTTP Server mode accepting container manifests over a simple API.

kube-proxy

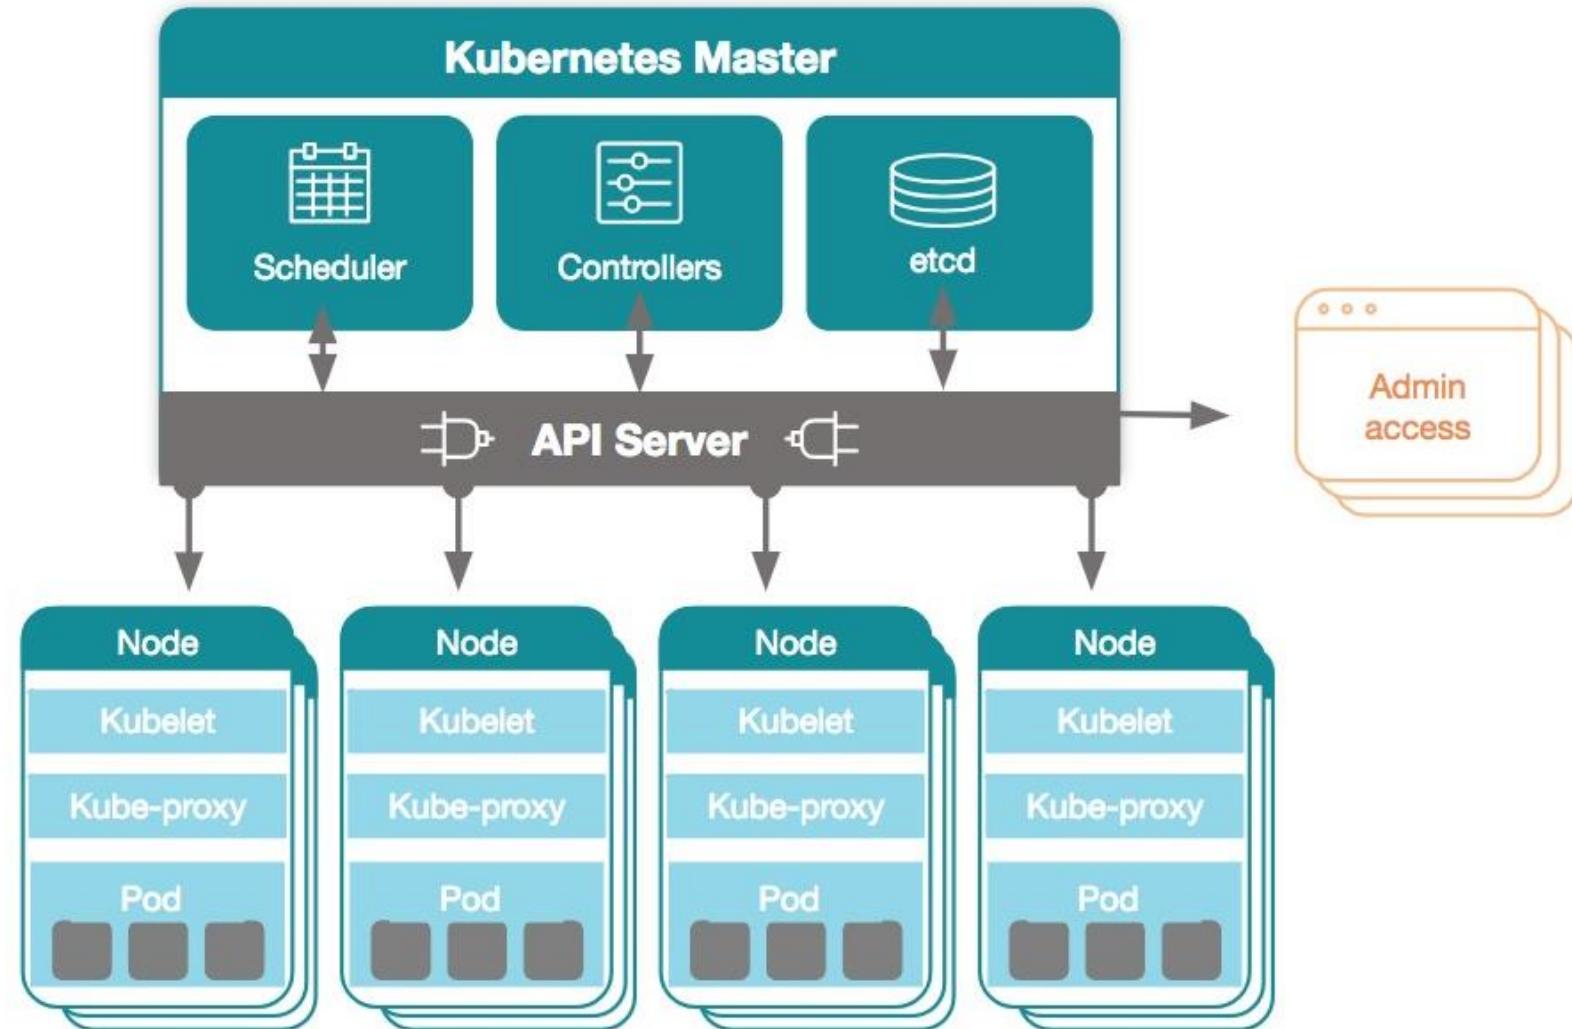
- Manages the network rules on each node.
- Performs connection forwarding or load balancing for Kubernetes cluster services.
- Available Proxy Modes:
 - Userspace
 - iptables
 - ipvs (default if supported)

Container Runtime Engine

- A container runtime is a CRI (Container Runtime Interface) compatible application that executes and manages containers.
 - Containerd (docker)
 - Cri-o
 - Rkt
 - Kata (formerly clear and hyper)
 - Virtlet (VM CRI compatible runtime)



Simplified Architectural View



Optional Services - Architecture Overview

- cloud-controller-manager
- Cluster DNS
- Kube Dashboard
- Heapster / Metrics API Server

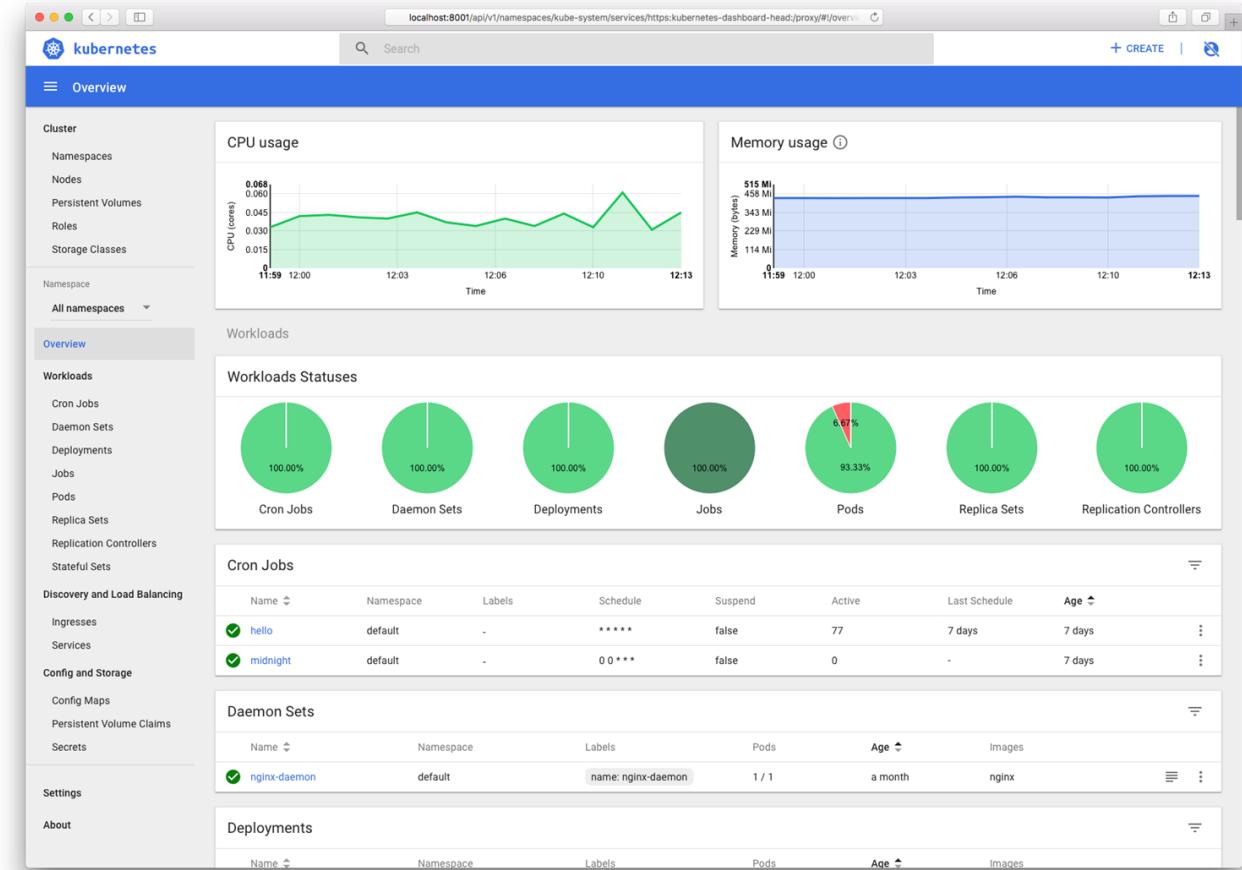
- Daemon that provides cloud-provider specific knowledge and integration capability into the core control loop of Kubernetes.
- The controllers include Node, Route, Service, and add an additional controller to handle things such as PersistentVolume Labels.

Cluster DNS

- Provides Cluster Wide DNS for Kubernetes Services.
 - Built on top of CoreDNS

Kube Dashboard

- A limited, general purpose web front end for the Kubernetes Cluster.



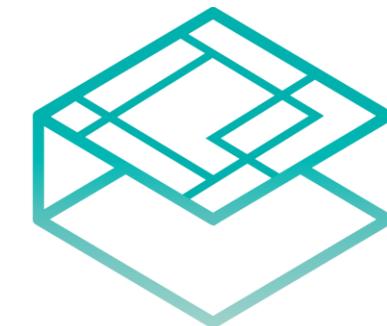
Metrics API Server

- Provides metrics for use with other Kubernetes Components.
 - Metrics API (current)

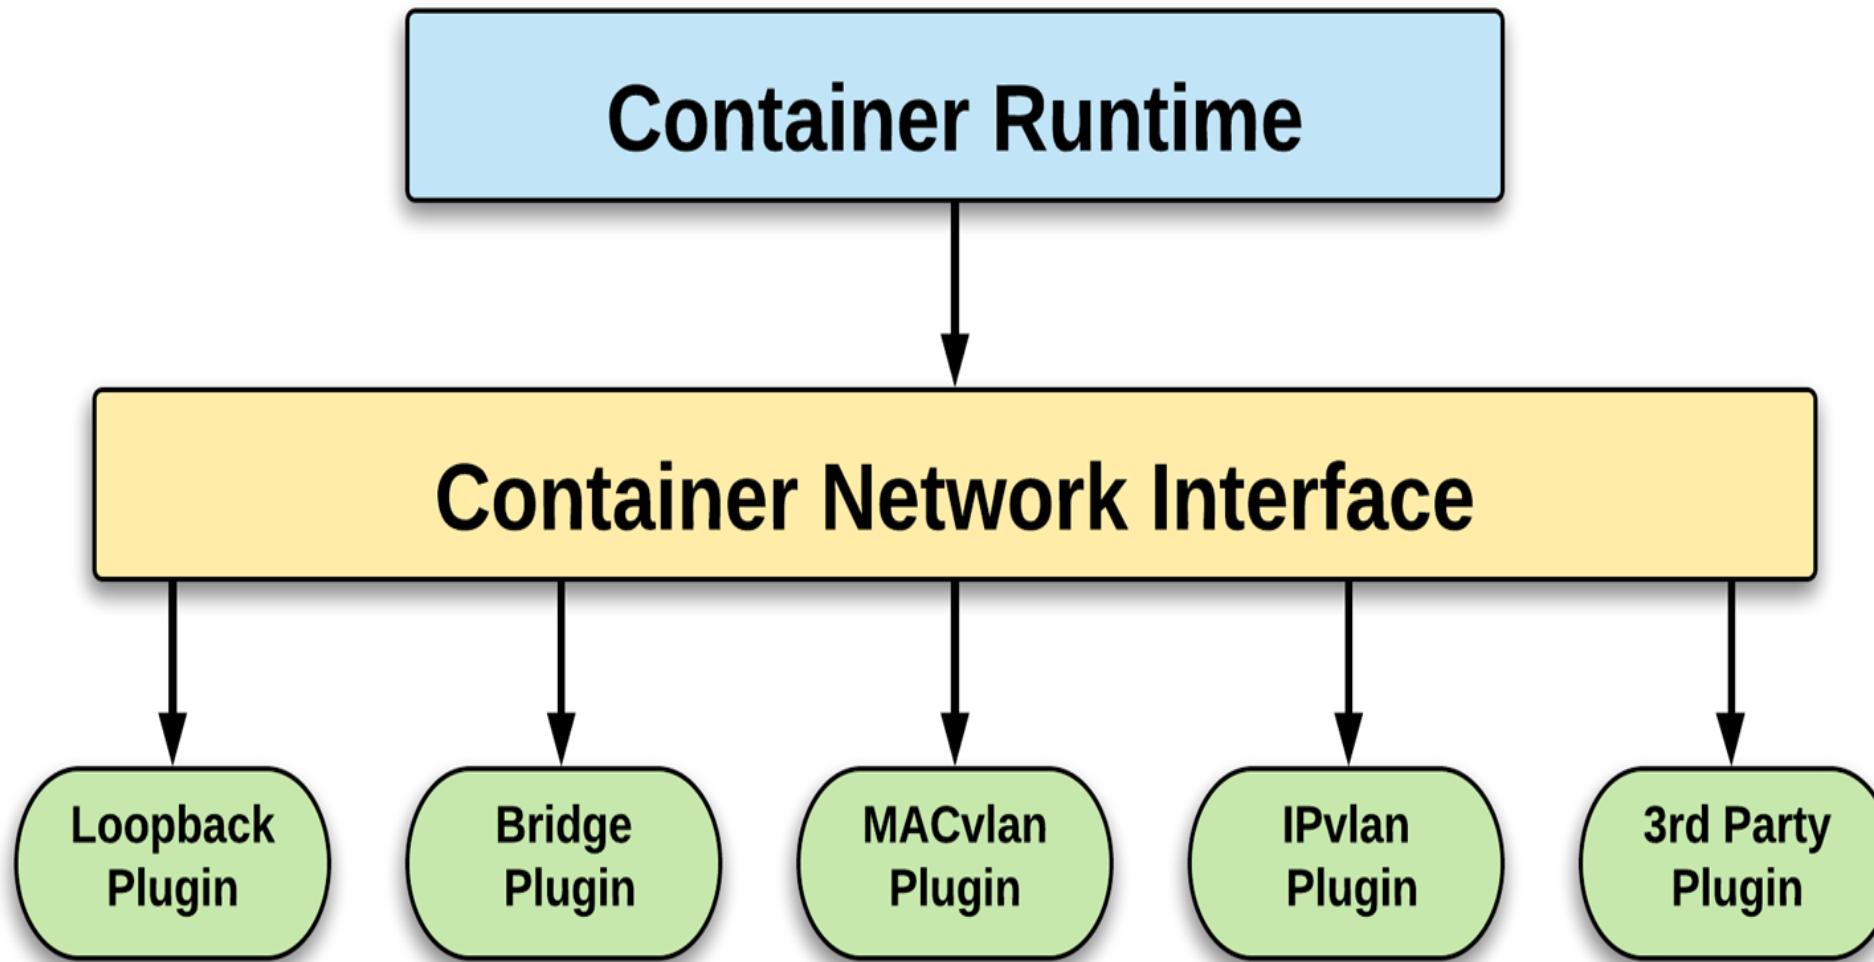
- **Pod Network**
 - Cluster-wide network used for pod-to-pod communication managed by a CNI (Container Network Interface) plugin.
- **Service Network**
 - Cluster-wide range of **Virtual IPs** managed by **kube-proxy** for service discovery.

Container Network Interface (CNI)

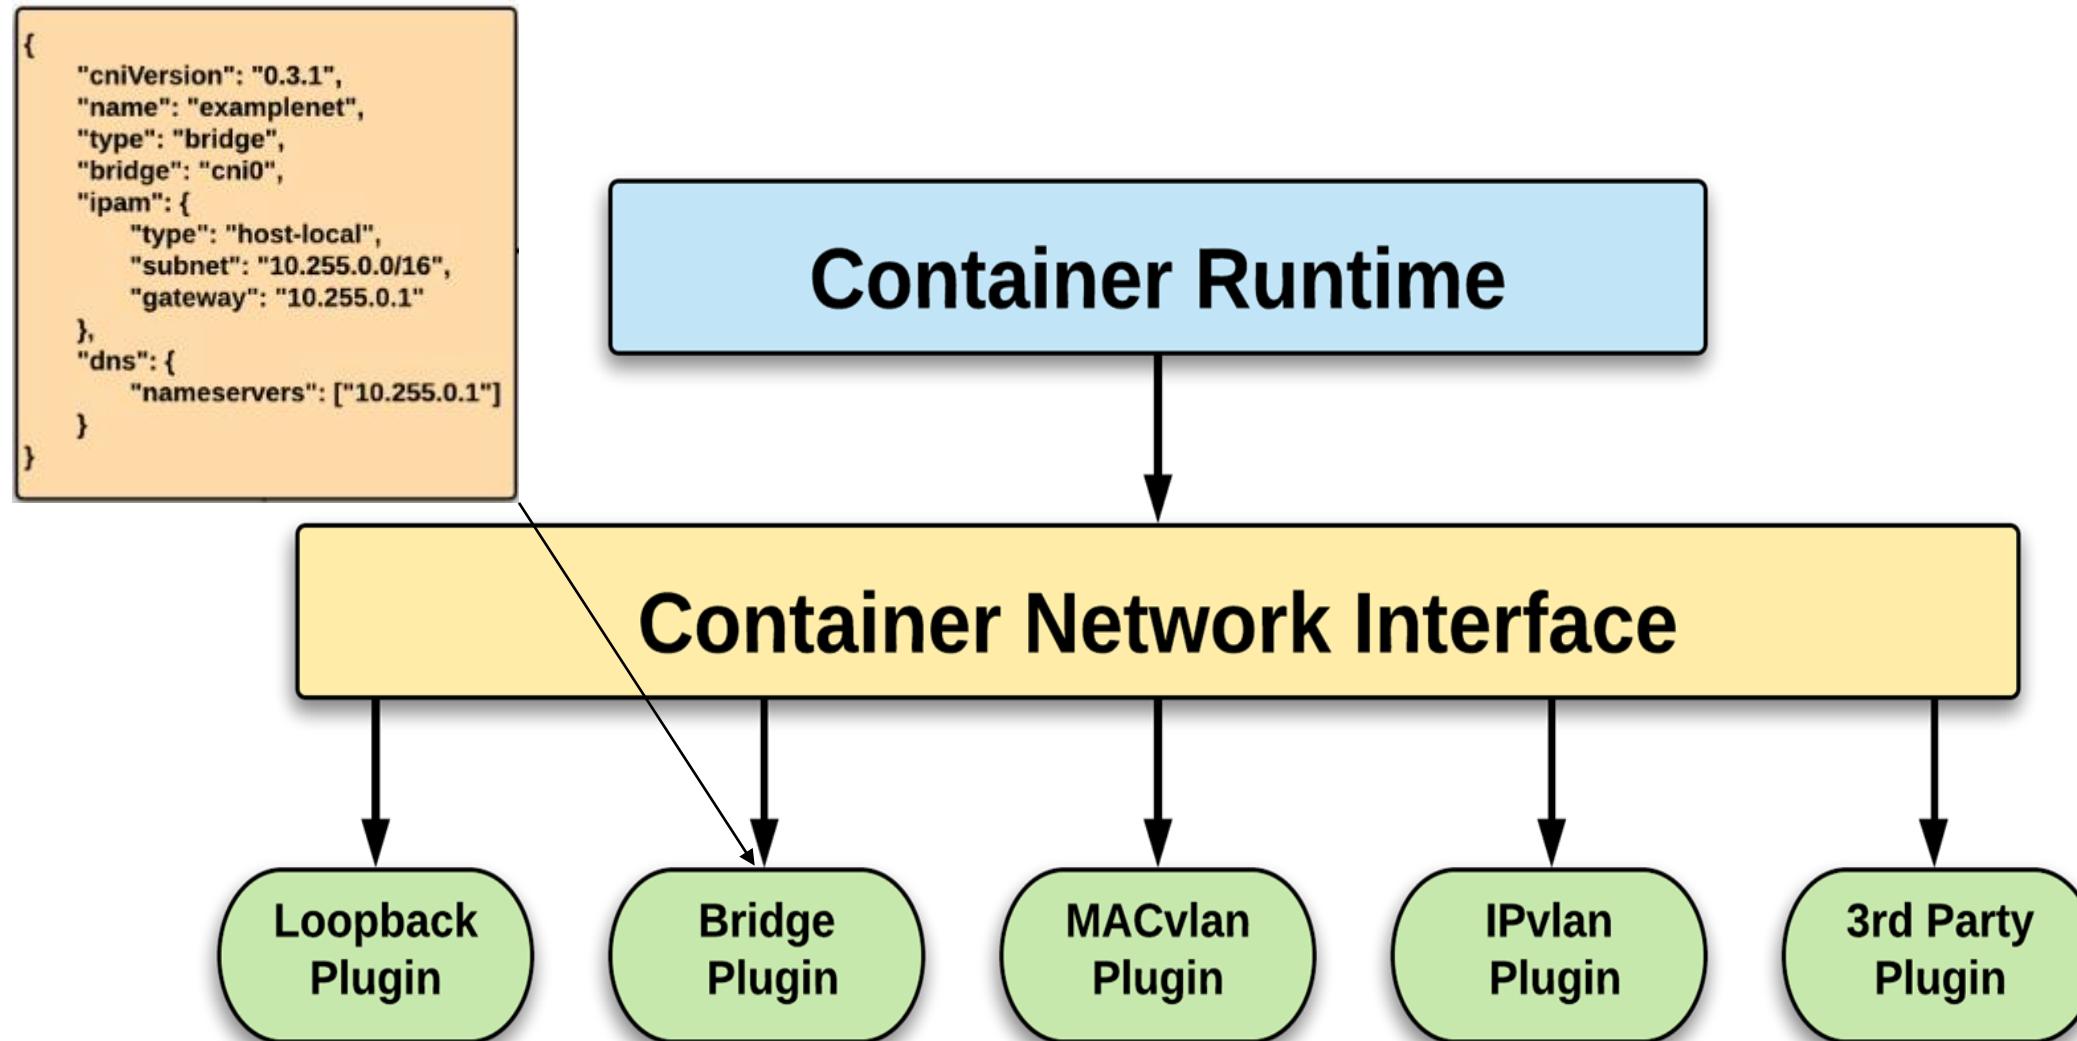
- Pod networking within Kubernetes is plumbed via the Container Network Interface (CNI).
- Functions as an interface between the container runtime and a **network implementation plugin**.
- CNCF Project
- Uses a simple JSON Schema.



C N I



CNI Overview

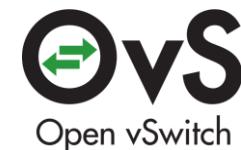


CNI Plugins

- Amazon ECS
- Calico
- Cilium
- Contiv
- Contrail
- Flannel



- GCE
- kube-router
- Multus
- OpenVSwitch
- Romana
- Weave



Fundamental Networking Rules

- All containers within a pod can communicate with each other unimpeded.
- All Pods can communicate with all other Pods without NAT.
- All nodes can communicate with all Pods (and vice-versa) without NAT.
- The IP that a Pod sees itself as is the same IP that others see it as.

- **Container-to-Container**

- Containers within a pod exist within the **same network namespace** and share an IP.
- Enables intrapod communication over localhost.

- **Pod-to-Pod**

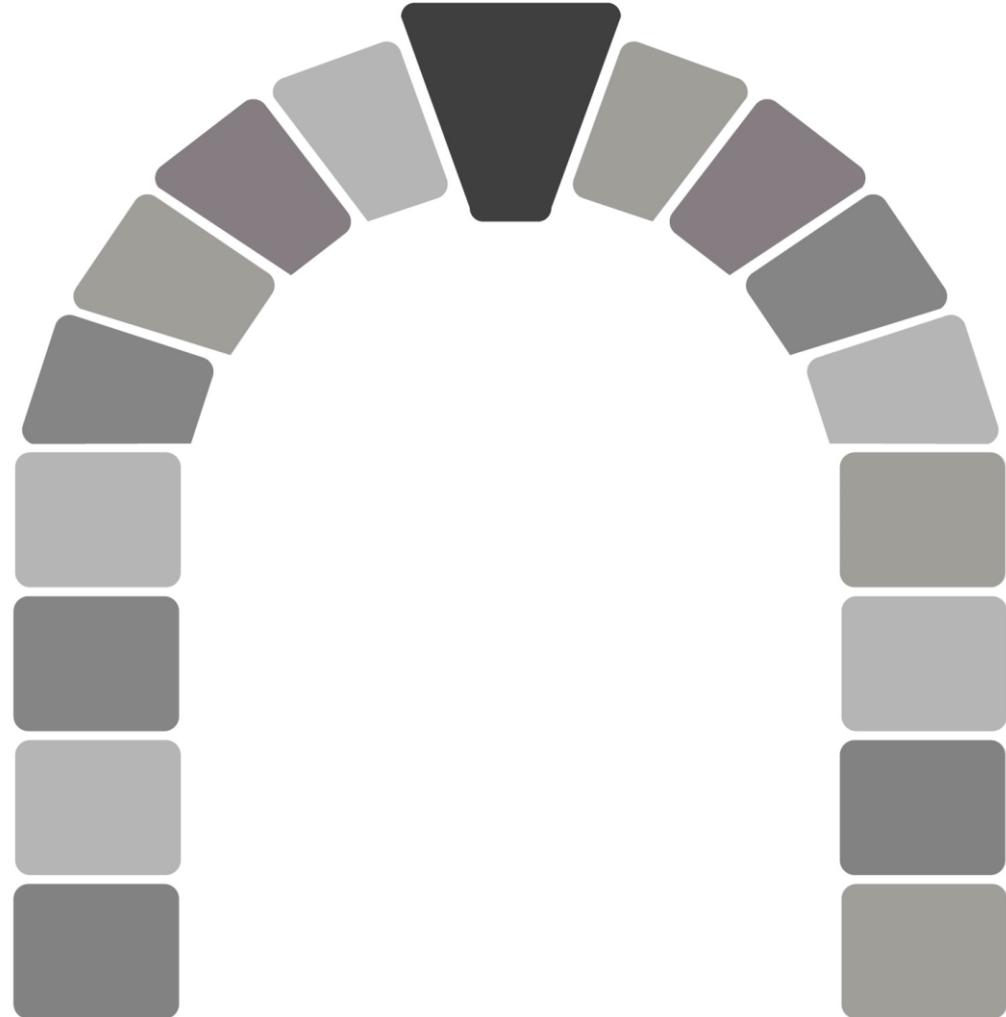
- Allocated **cluster unique IP** for the duration of its life cycle.
- Pods themselves are fundamentally ephemeral.

- **Pod-to-Service**
 - managed by **kube-proxy** and given a **persistent cluster unique IP**
 - exists beyond a Pod's lifecycle.
- **External-to-Service**
 - Handled by **kube-proxy**.
 - Works in cooperation with a cloud provider or other external entity (load balancer).

Concepts and Resources

The API and Object Model

- The **REST API** is the true **keystone** of Kubernetes.
- **Everything** within the Kubernetes is as an **API Object**.



API Groups

- Designed to make it extremely simple to both understand and extend.
- An API Group is a **REST compatible path** that acts as the type descriptor for a Kubernetes object.
- Referenced within an object as the **apiVersion** and **kind**.

Format:

`/apis/<group>/<version>/<resource>`

Examples:

`/apis/apps/v1/deployments`

`/apis/batch/v1beta1/cronjobs`

API Versioning

- Three tiers of API maturity levels.
- Also referenced within the object **apiVersion**.

Format:

`/apis/<group>/<version>/<resource>`

Examples:

`/apis/apps/v1/deployments`

`/apis/batch/v1beta1/cronjobs`

- **Alpha:** Possibly buggy, And may change. **Disabled by default.**
- **Beta:** Tested and considered stable. However API Schema may change. **Enabled by default.**
- **Stable:** Released, stable and API schema will not change. **Enabled by default.**

Object Model

- Objects are a “record of intent” or a persistent entity that represent the desired state of the object within the cluster.
- All objects **MUST** have `apiVersion`, `kind`, and possess the nested fields `metadata.name`, `metadata.namespace`, and `metadata.uid`.

Object Model Requirements

- **apiVersion**: Kubernetes API version of the Object
- **kind**: Type of Kubernetes Object
- **metadata.name**: Unique name of the Object
- **metadata.namespace**: Scoped environment name that the object belongs to (will default to current).
- **metadata.uid**: The (generated) uid for an object.

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-example
  namespace: default
  uid: f8798d82-1185-11e8-94ce-080027b3c7a6
```

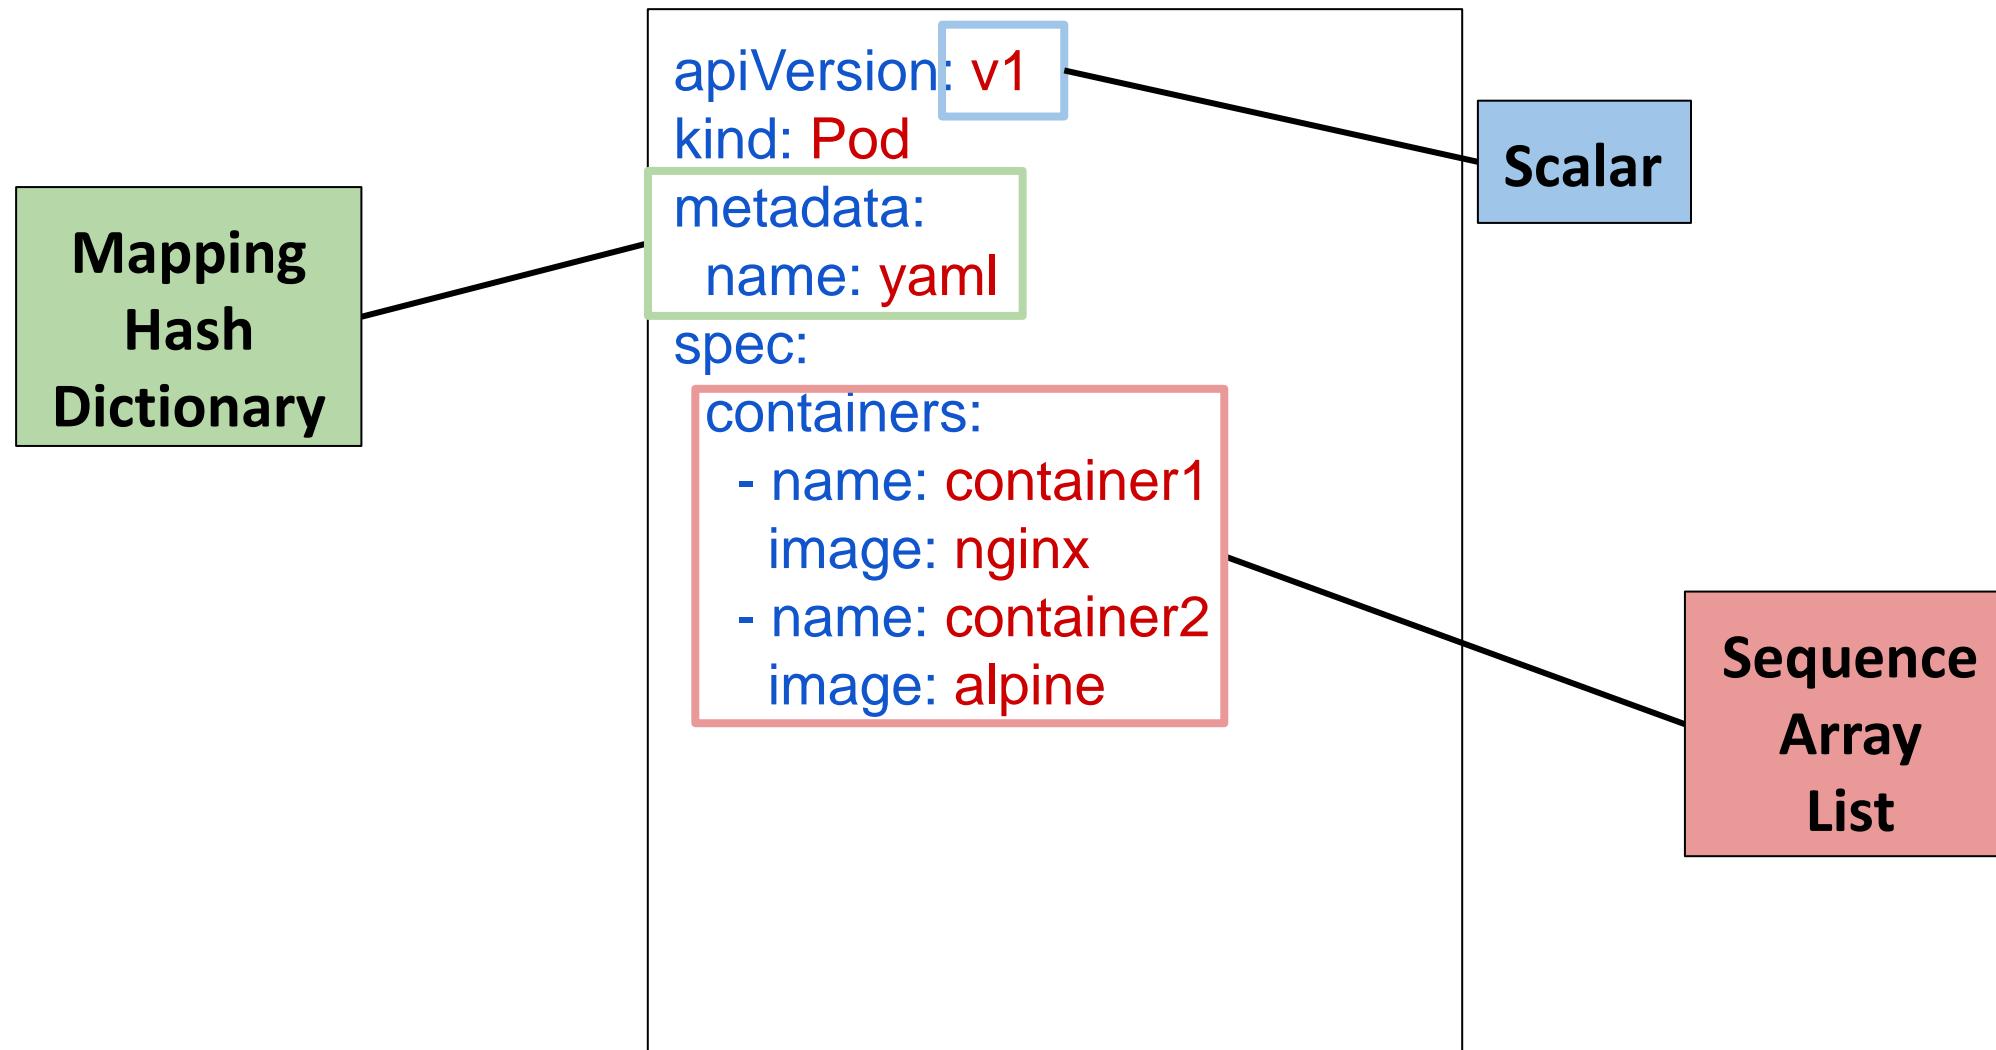
Object Expression - YAML

- Files or other representations of Kubernetes Objects are generally represented in YAML.
- A “Human Friendly” data serialization standard.
- Uses white space (specifically spaces) alignment to denote ownership.
- Three basic data types:
 - **mappings** - hash or dictionary,
 - **sequences** - array or list
 - **scalars** - string, number, boolean etc

Object Expression - YAML

```
apiVersion: v1
kind: Pod
metadata:
  name: yaml
spec:
  containers:
    - name: container1
      image: nginx
    - name: container2
      image: alpine
```

Object Expression - YAML



YAML vs JSON

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-example
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
    ports:
      - containerPort: 80
```

```
{
  "apiVersion": "v1",
  "kind": "Pod",
  "metadata": {
    "name": "pod-example"
  },
  "spec": {
    "containers": [
      {
        "name": "nginx",
        "image": "nginx:stable-alpine",
        "ports": [ { "containerPort": 80 } ]
      }
    ]
  }
}
```

Object Model - Workloads

- Workload related objects within Kubernetes have an additional two nested fields **spec** and **status**.
 - **spec** - Describes the **desired state or configuration** of the object to be created.
 - **status** - Is managed by Kubernetes and describes the **actual state** of the object and its history.

Workload Object Example

Example Object

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-example
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
      ports:
        - containerPort: 80
```

Example Status Snippet

```
status:
  conditions:
    - lastProbeTime: null
      lastTransitionTime: 2018-02-14T14:15:52Z
      status: "True"
      type: Ready
    - lastProbeTime: null
      lastTransitionTime: 2018-02-14T14:15:49Z
      status: "True"
      type: Initialized
    - lastProbeTime: null
      lastTransitionTime: 2018-02-14T14:15:49Z
      status: "True"
      type: PodScheduled
```

Core Objects - Concepts and Resources

- **Namespaces**
- **Pods**
- **Labels**
- **Selectors**
- **Services**

- Kubernetes has several core building blocks that make up the foundation of their higher level components.

Namespaces

Pods

Labels

Services

Selectors

Namespaces

- Namespaces are a logical cluster or environment, and are the primary method of partitioning a cluster or scoping access.

```
apiVersion: v1
kind: Namespace
metadata:
  name: prod
  labels:
    app: MyBigWebApp
```

```
$ kubectl get ns --show-labels
NAME      STATUS  AGE   LABELS
default   Active  11h   <none>
kube-public   Active  11h   <none>
kube-system   Active  11h   <none>
prod       Active  6s    app=MyBigWebApp
```

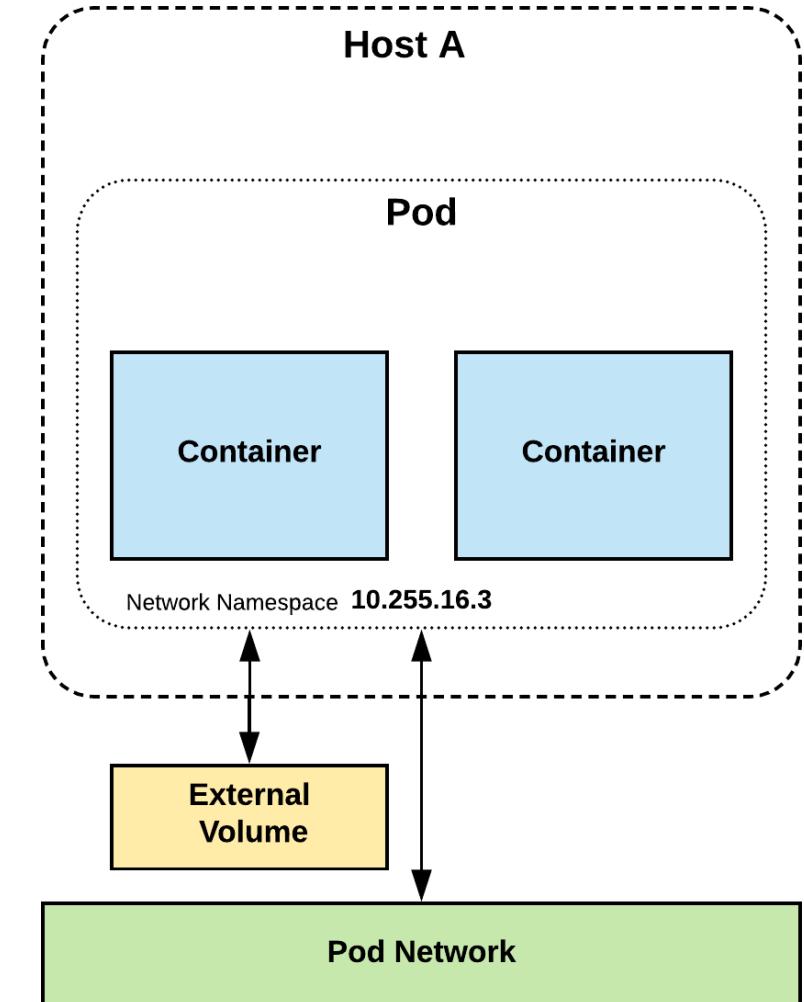
Default Namespaces

- **default**: The default namespace for any object without a namespace.
- **kube-system**: Acts as the home for objects and resources created by Kubernetes itself.
- **kube-public**: A special namespace; readable by all users that is reserved for cluster bootstrapping and configuration.

```
$ kubectl get ns --show-labels
NAME      STATUS  AGE    LABELS
default   Active  11h   <none>
kube-public Active  11h   <none>
kube-system Active  11h   <none>
```

Pod

- **Atomic unit** or smallest “unit of work” of Kubernetes.
- Foundational building block of Kubernetes Workloads.
- Pods are one or more containers that share volumes, a network namespace, and are a part of a **single context**.



Pod Examples

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-example
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
  ports:
    - containerPort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: multi-container-example
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
      ports:
        - containerPort: 80
      volumeMounts:
        - name: html
          mountPath: /usr/share/nginx/html
    - name: content
      image: alpine:latest
      command: ["/bin/sh", "-c"]
      args:
        - while true; do
            date >> /html/index.html;
            sleep 5;
        done
      volumeMounts:
        - name: html
          mountPath: /html
  volumes:
    - name: html
      emptyDir: {}
```

Key Pod Container Attributes

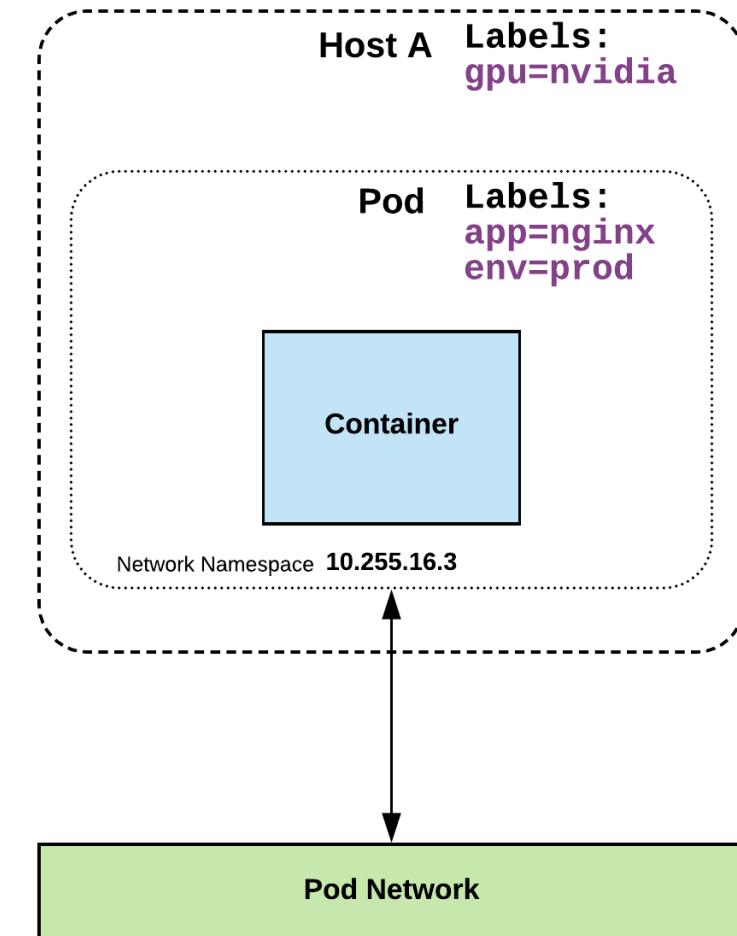
- **name** - The name of the container
- **image** - The container image
- **ports** - array of ports to expose. Can be granted a friendly name and protocol may be specified
- **env** - array of environment variables
- **command** - Entrypoint array (equiv to Docker **ENTRYPOINT**)
- **args** - Arguments to pass to the command (equiv to Docker **CMD**)

Container

```
name: nginx
image: nginx:stable-alpine
ports:
- containerPort: 80
  name: http
  protocol: TCP
env:
- name: MYVAR
  value: isAwesome
command: ["/bin/sh", "-c"]
args: ["echo ${MYVAR}"]
```

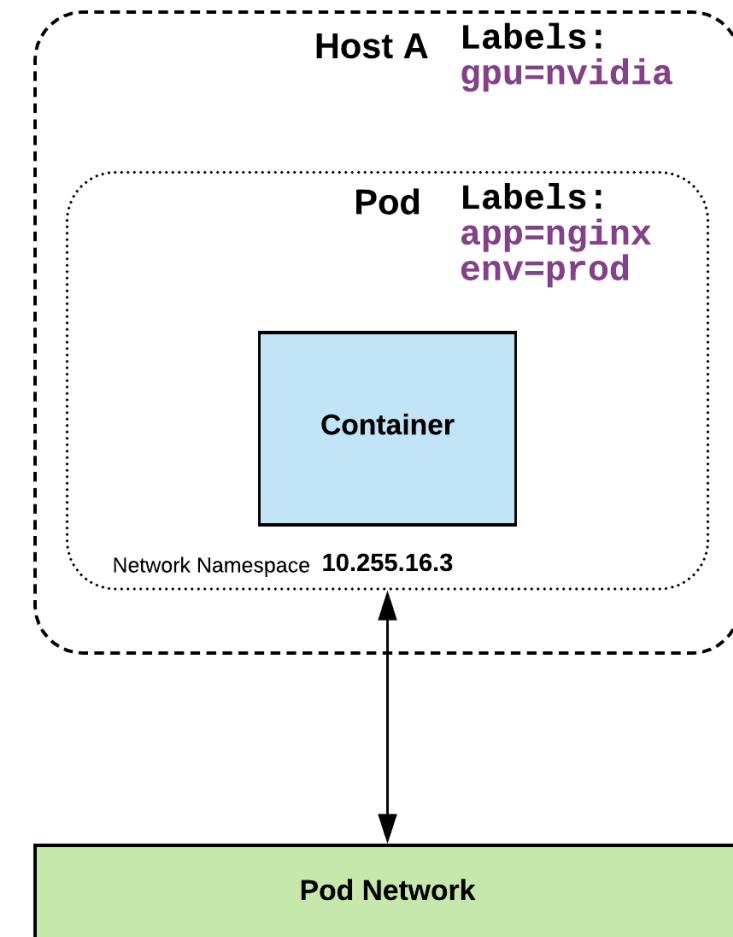
Labels

- key-value pairs that are used to identify, describe and group together related sets of objects or resources.
- **NOT** characteristic of uniqueness.
- Have a strict syntax with a slightly limited character set*.



Label Example

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-label-example
  labels:
    app: nginx
    env: prod
spec:
  containers:
  - name: nginx
    image: nginx:stable-alpine
    ports:
    - containerPort: 80
```



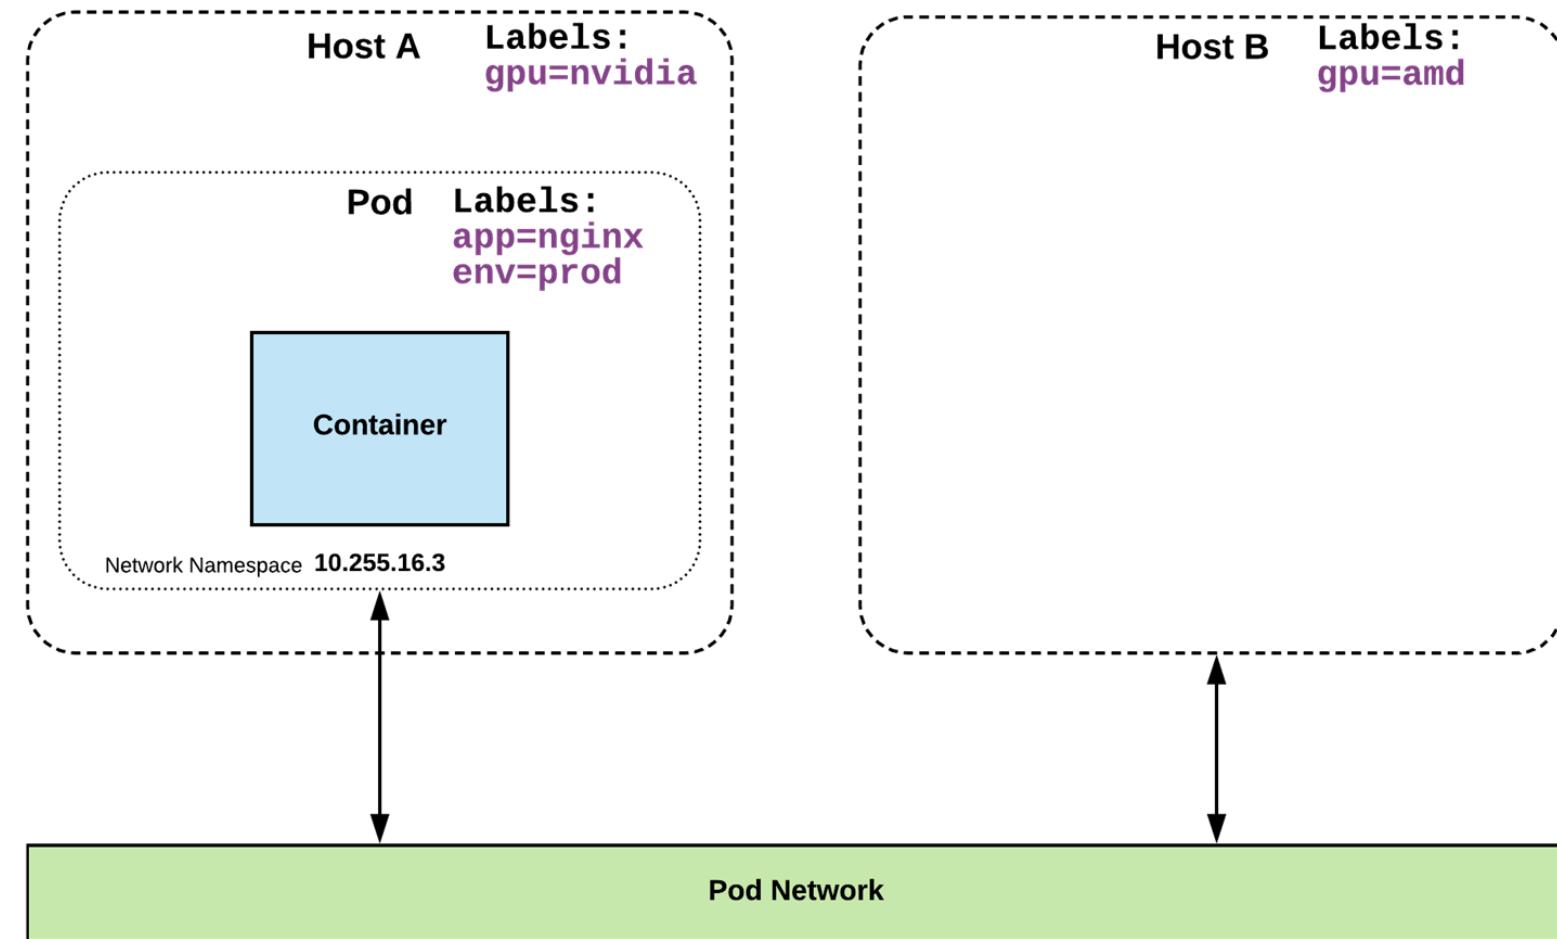
Selectors

- Selectors use labels to filter or select objects, and are used throughout Kubernetes.

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-label-example
labels:
  app: nginx
  env: prod
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
      ports:
        - containerPort: 80
nodeSelector:
  gpu: nvidia
```

Selectors Example

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-label-example
  labels:
    app: nginx
    env: prod
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
      ports:
        - containerPort: 80
nodeSelector:
  gpu: nvidia
```



Selectors Types

Equality based selectors allow for simple filtering (=, ==, or !=).

```
selector:  
matchLabels:  
gpu: nvidia
```

Set-based selectors are supported on a limited subset of objects. However, they provide a method of filtering on a set of values, and supports multiple operators including: **in**, **notin**, and **exist**.

```
selector:  
matchExpressions:  
- key: gpu  
operator: in  
values: ["nvidia"]
```

- **Unified method of accessing** the exposed workloads of Pods.
- **Durable resource** (unlike Pods)
 - static cluster-unique IP
 - static namespaced DNS name

<service name>.<namespace>.svc.cluster.local

- Target Pods using **equality based selectors**.
- Uses **kube-proxy** to provide simple load-balancing.
- **kube-proxy** acts as a daemon that creates **local entries** in the host's iptables for every service.

Service Types

- There are 4 major service types:

- ClusterIP (default)
- NodePort
- LoadBalancer
- ExternalName

ClusterIP Service

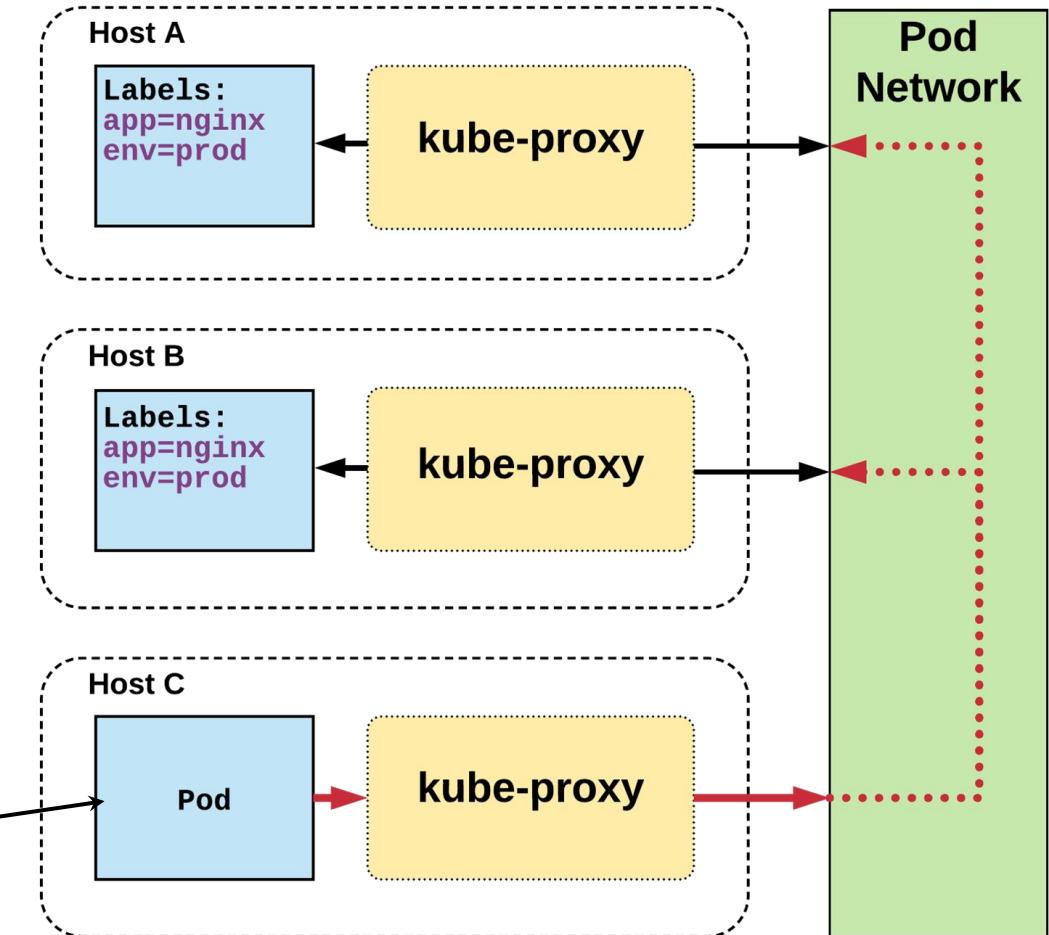
- **ClusterIP** services exposes a service on a strictly cluster internal virtual IP.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  selector:
    app: nginx
    env: prod
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

Cluster IP Service

```
Name: example-prod
Selector: app=nginx,env=prod
Type: ClusterIP
IP: 10.96.28.176
Port: <unset> 80/TCP
TargetPort: 80/TCP
Endpoints: 10.255.16.3:80,
           10.255.16.4:80
```

```
/ # nslookup example-prod.default.svc.cluster.local
Name: example-prod.default.svc.cluster.local
Address 1: 10.96.28.176 example-prod.default.svc.cluster.local
```

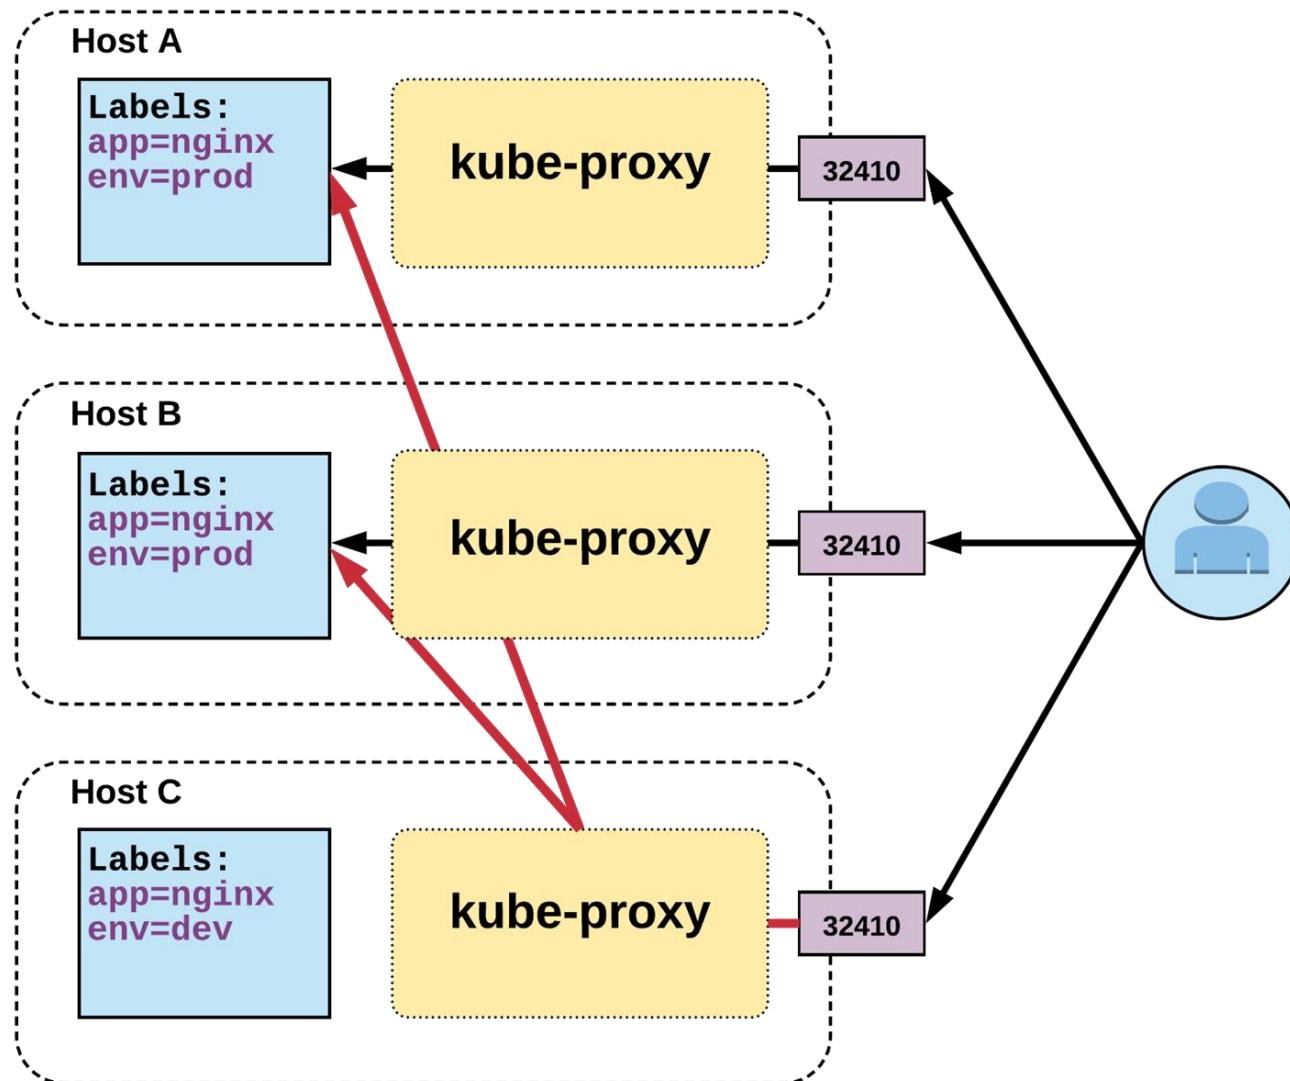


NodePort Service

- **NodePort** services extend the **ClusterIP** service.
- Exposes a port on every node's IP.
- Port can either be statically defined, or dynamically taken from a range between 30000-32767.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: NodePort
  selector:
    app: nginx
    env: prod
  ports:
    - nodePort: 32410
      protocol: TCP
      port: 80
      targetPort: 80
```

NodePort Service



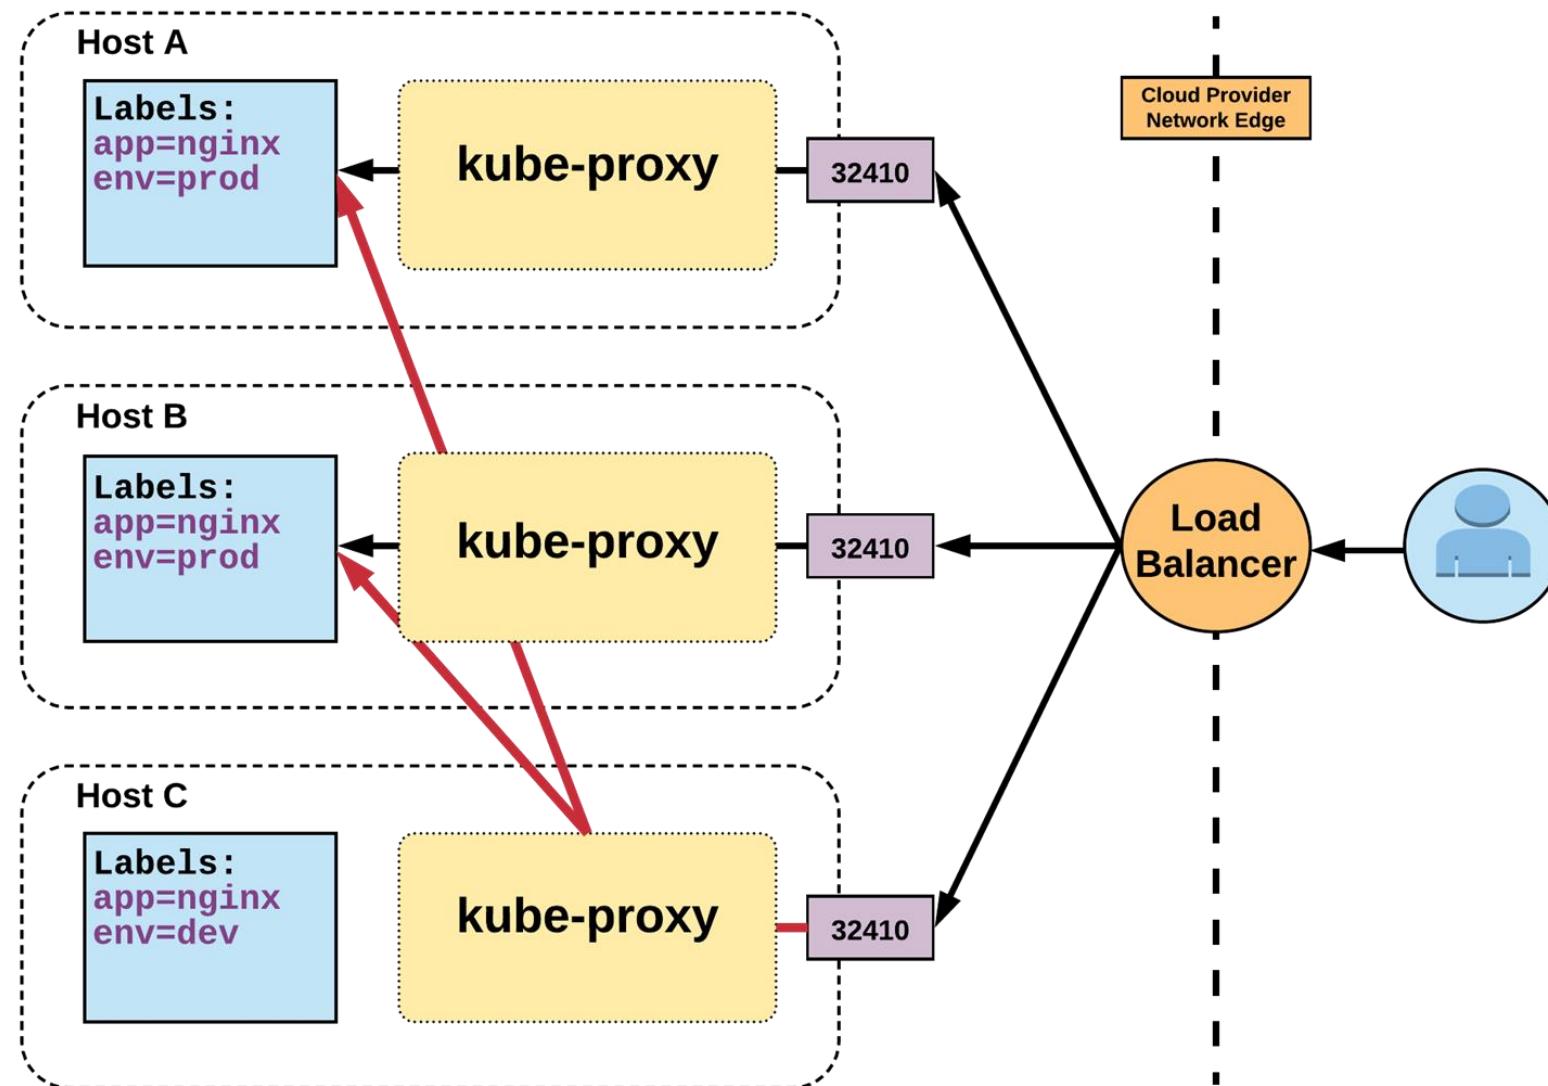
Name:	example-prod
Selector:	app=nginx,env=prod
Type:	NodePort
IP:	10.96.28.176
Port:	<unset> 80/TCP
TargetPort:	80/TCP
NodePort:	<unset> 32410/TCP
Endpoints:	10.255.16.3:80, 10.255.16.4:80

LoadBalancer Service

- **LoadBalancer** services extend **NodePort**.
- Works in conjunction with an external system to map a cluster external IP to the exposed service.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: LoadBalancer
  selector:
    app: nginx
    env: prod
  ports:
    protocol: TCP
    port: 80
    targetPort: 80
```

LoadBalancer Service



Name:	example-prod
Selector:	app=nginx,env=prod
Type:	LoadBalancer
IP:	10.96.28.176
LoadBalancer	
Ingress:	172.17.18.43
Port:	<unset> 80/TCP
TargetPort:	80/TCP
NodePort:	<unset> 32410/TCP
Endpoints:	10.255.16.3:80, 10.255.16.4:80

ExternalName Service

- **ExternalName** is used to reference endpoints **OUTSIDE** the cluster.
- Creates an internal **CNAME** DNS entry that aliases another.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: ExternalName
spec:
  externalName: example.com
```

Exploring the Core

Workloads - Concepts and Resources

- **ReplicaSet**
- **Deployment**
- **DaemonSet**
- **StatefulSet**
- **Job**
- **CronJob**

Workloads

- Workloads within Kubernetes are higher level objects that manage Pods or other higher level objects.
- In **ALL CASES** a Pod Template is included, and acts the base tier of management.

Pod Template

- Workload Controllers manage instances of Pods based off a provided template.
- Pod Templates are Pod specs with limited metadata.
- Controllers use Pod Templates to make actual pods.

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-example
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx
```

```
template:
  metadata:
    labels:
      app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx
```

ReplicaSet

- Primary method of managing pod replicas and their lifecycle.
- Includes their scheduling, scaling, and deletion.
- Their job is simple: **Always ensure the desired number of pods are running.**



ReplicaSet

- **replicas**: The desired number of instances of the Pod.
- **selector**: The label selector for the **ReplicaSet** will manage **ALL** Pod instances that it targets; whether it's desired or not.

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rs-example
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
      env: prod
  template:
    <pod template>
```

ReplicaSet

```

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rs-example
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
      env: prod
  template:
    metadata:
      labels:
        app: nginx
        env: prod
    spec:
      containers:
        - name: nginx
          image: nginx:stable-alpine
      ports:
        - containerPort: 80
  
```

```

$ kubectl get pods
NAME           READY   STATUS  RESTARTS AGE
rs-example-9l4dt  1/1    Running 0       1h
rs-example-b7bcg  1/1    Running 0       1h
rs-example-mkll2  1/1    Running 0       1h
  
```

```

$ kubectl describe rs rs-example
Name: rs-example
Namespace: default
Selector: app=nginx,env=prod
Labels: app=nginx
        env=prod
Annotations: <none>
Replicas: 3 current / 3 desired
Pods Status: 3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels: app=nginx
          env=prod
  Containers:
    nginx:
      Image: nginx:stable-alpine
      Port: 80/TCP
      Environment: <none>
      Mounts: <none>
      Volumes: <none>
  Events:
    Type  Reason     Age From           Message
    ----  -----    --  ----           -----
    Normal SuccessfulCreate 16s replicaset-controller Created pod: rs-example-mkll2
    Normal SuccessfulCreate 16s replicaset-controller Created pod: rs-example-b7bcg
    Normal SuccessfulCreate 16s replicaset-controller Created pod: rs-example-9l4d
  
```

Deployment

- Declarative method of managing Pods via ReplicaSets.
- Provide rollback functionality and update control.
- Updates are managed through the pod-template-hash label.
- Each iteration creates a unique label that is assigned to both the ReplicaSet and subsequent Pods.



Deployment

- **revisionHistoryLimit**: The number of previous iterations of the Deployment to retain.
- **strategy**: Describes the method of updating the Pods based on the **type**. Valid options are **Recreate** or **RollingUpdate**.
 - **Recreate**: All existing Pods are killed before the new ones are created.
 - **RollingUpdate**: Cycles through updating the Pods according to the parameters: **maxSurge** and **maxUnavailable**.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: deploy-example
spec:
  replicas: 3
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
      env: prod
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  template:
    <pod template>
```

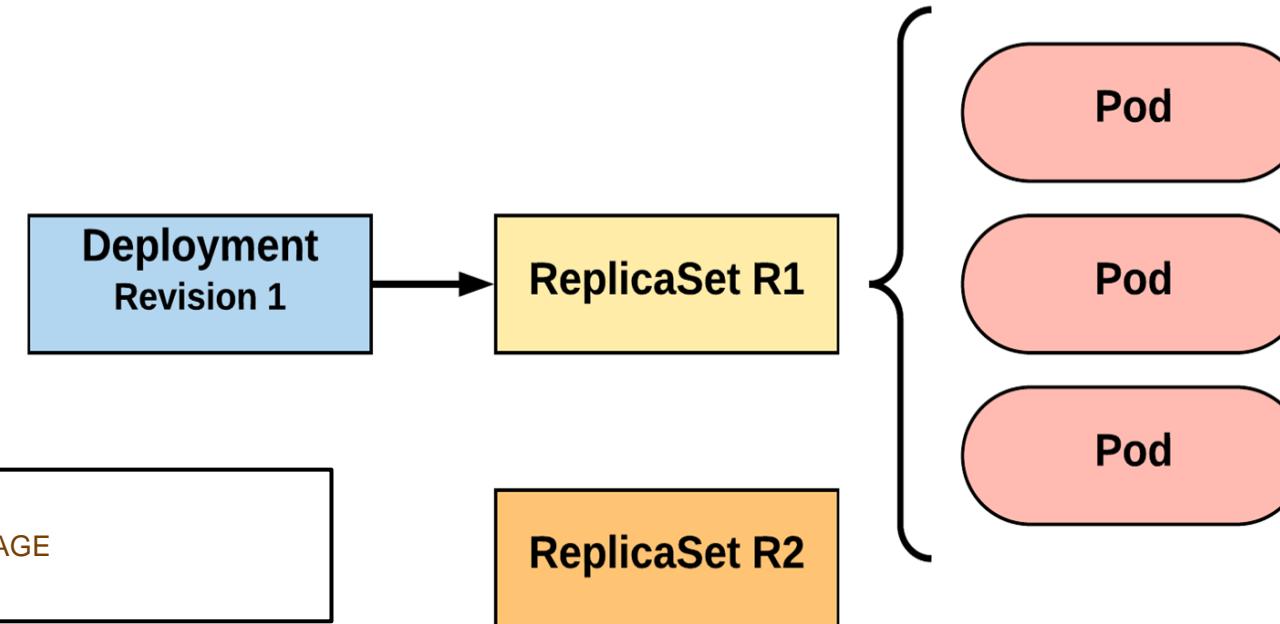
RollingUpdate Deployment

- Updating pod template generates a new **ReplicaSet** revision.

```
R1 pod-template-hash:  
676677fff  
R2 pod-template-hash:  
54f7ff7d6d
```

```
$ kubectl get replicaset  
NAME      DESIRED  CURRENT  READY   AGE  
mydep-676677fff  3        3        3      5h
```

```
$ kubectl get pods  
NAME      READY  STATUS  RESTARTS  AGE  
mydep-676677fff-9r2zn  1/1    Running  0      5h  
mydep-676677fff-hsfz9  1/1    Running  0      5h  
mydep-676677fff-sjxhf  1/1    Running  0      5h
```



RollingUpdate Deployment Types

- New ReplicaSet is initially scaled up based on **maxSurge**.

```
R1 pod-template-hash:
```

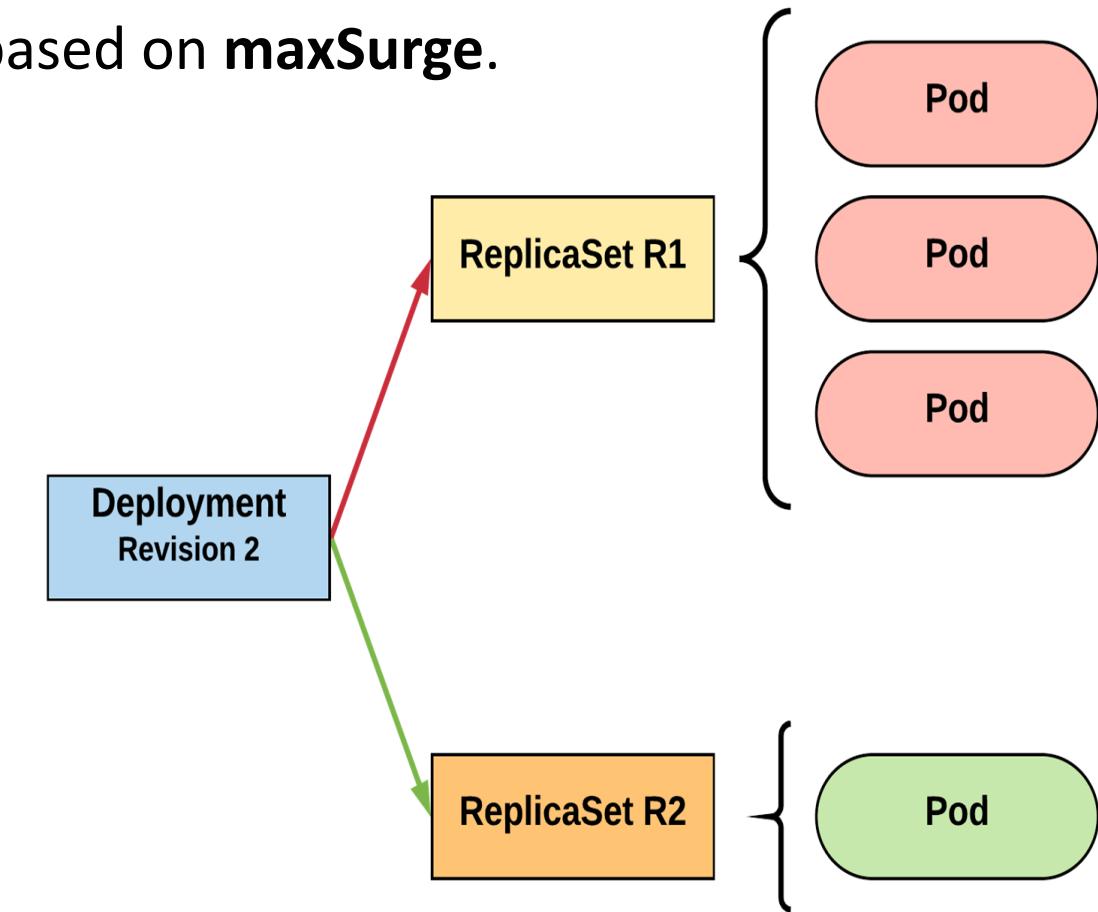
```
676677fff
```

```
R2 pod-template-hash:
```

```
54f7ff7d6d
```

```
$ kubectl get replicaset  
NAME      DESIRED  CURRENT  READY   AGE  
mydep-54f7ff7d6d  1        1        1      5s  
mydep-6766777fff  2        3        3      5h
```

```
$ kubectl get pods  
NAME      READY  STATUS  RESTARTS  AGE  
mydep-54f7ff7d6d-9gvll  1/1  Running  0      2s  
mydep-6766777fff-9r2zn  1/1  Running  0      5h  
mydep-6766777fff-hsfz9  1/1  Running  0      5h  
mydep-6766777fff-sjxhf  1/1  Running  0      5h
```



RollingUpdate Deployment

- Phase out of old Pods managed by **maxSurge** and **maxUnavailable**.

R1 pod-template-hash:

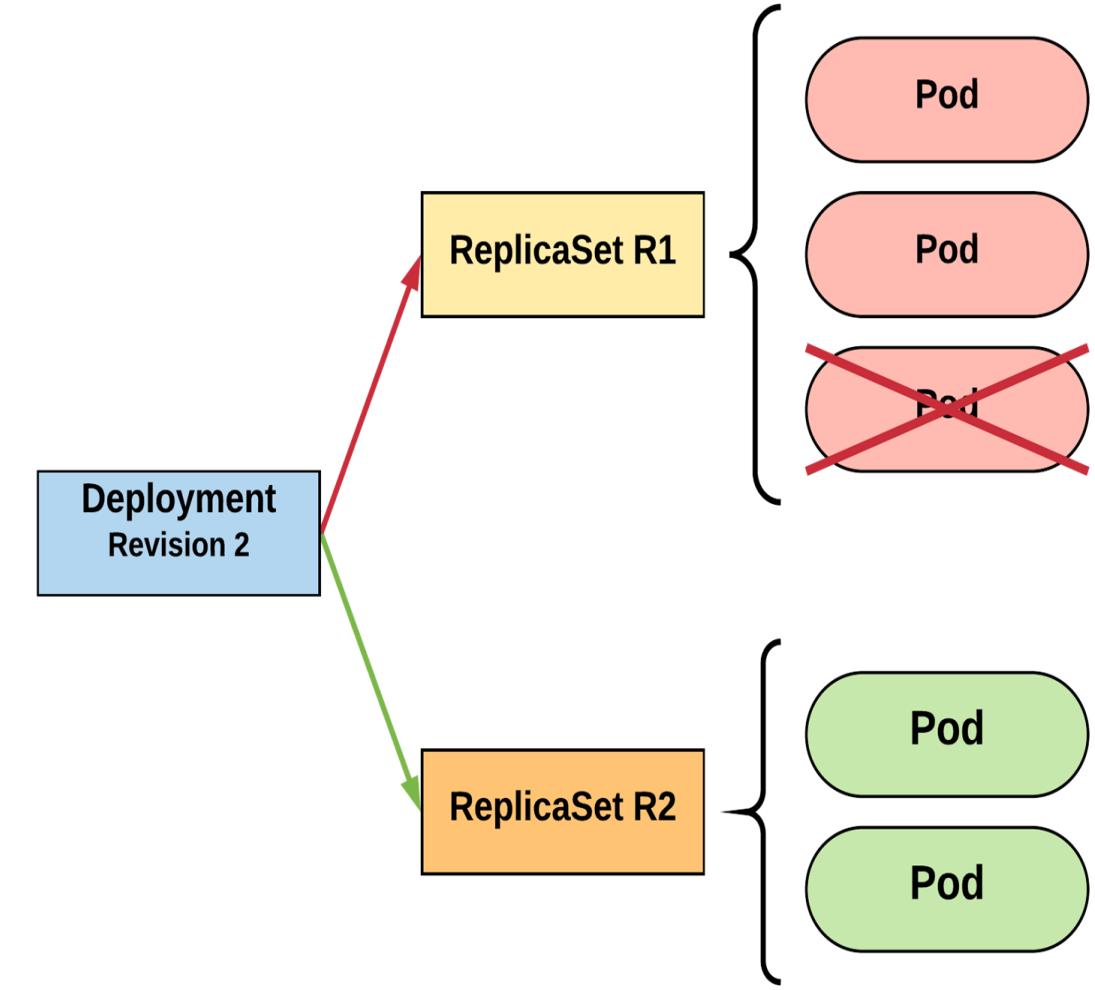
676677fff

R2 pod-template-hash:

54f7ff7d6d

```
$ kubectl get replicaset
NAME      DESIRED  CURRENT  READY   AGE
mydep-54f7ff7d6d  2        2        2     8s
mydep-6766777fff  2        2        2     5h
```

```
$ kubectl get pods
NAME      READY  STATUS  RESTARTS  AGE
mydep-54f7ff7d6d-9gvll  1/1  Running  0       5s
mydep-54f7ff7d6d-cqvlq  1/1  Running  0       2s
mydep-6766777fff-9r2zn  1/1  Running  0       5h
mydep-6766777fff-hsfz9  1/1  Running  0       5h
```



RollingUpdate Deployment

- Phase out of old Pods managed by **maxSurge** and **maxUnavailable**.

R1 pod-template-hash:

676677fff

R2 pod-template-hash:

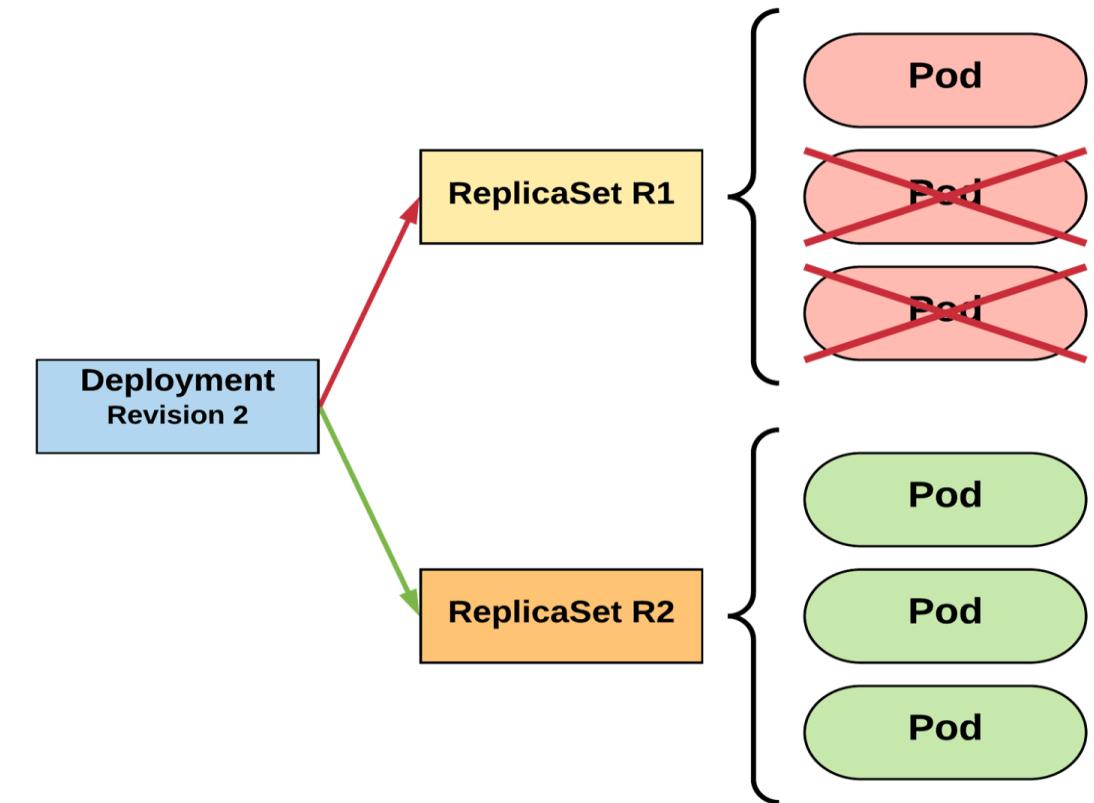
54f7ff7d6d

\$ kubectl get replicaset

NAME	DESIRED	CURRENT	READY	AGE
mydep-54f7ff7d6d	3	3	3	10s
mydep-676677fff	0	1	1	5h

\$ kubectl get pods

NAME	READY	STATUS	RESTARTS	AGE
mydep-54f7ff7d6d-9gvll	1/1	Running	0	7s
mydep-54f7ff7d6d-cqvlq	1/1	Running	0	5s
mydep-54f7ff7d6d-gccr6	1/1	Running	0	2s
mydep-676677fff-9r2zn	1/1	Running	0	5h



RollingUpdate Deployment

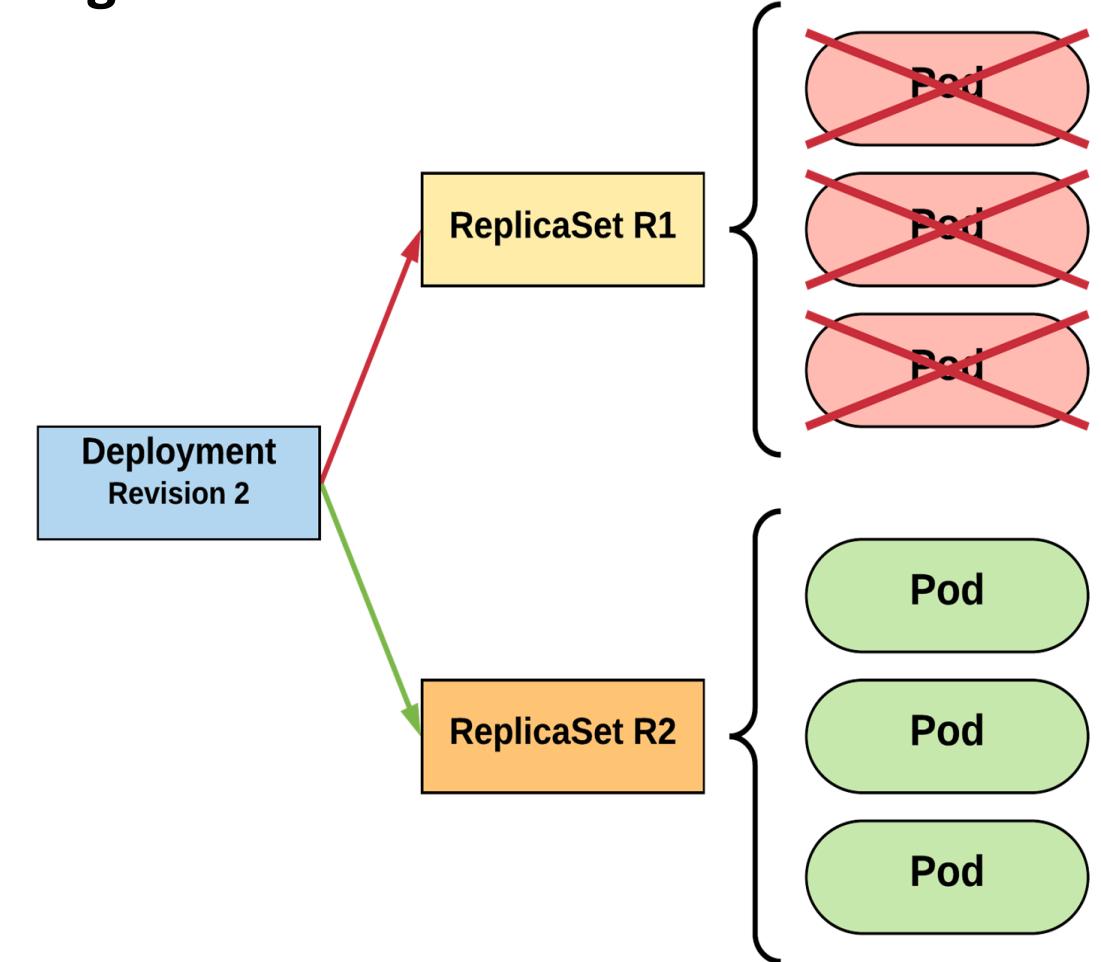
- Phase out of old Pods managed by **maxSurge** and **maxUnavailable**.

R1 pod-template-hash:
676677fff

R2 pod-template-hash:
54f7ff7d6d

```
$ kubectl get replicaset
NAME      DESIRED  CURRENT  READY   AGE
mydep-54f7ff7d6d  3        3        3     13s
mydep-676677fff  0        0        0     5h
```

```
$ kubectl get pods
NAME          READY  STATUS    RESTARTS  AGE
mydep-54f7ff7d6d-9gvll  1/1   Running   0         10s
mydep-54f7ff7d6d-cqvlq  1/1   Running   0         8s
mydep-54f7ff7d6d-gccr6  1/1   Running   0         5s
```



RollingUpdate Deployment

- Updated to new deployment revision completed.

R1 pod-template-hash:

676677fff

R2 pod-template-hash:

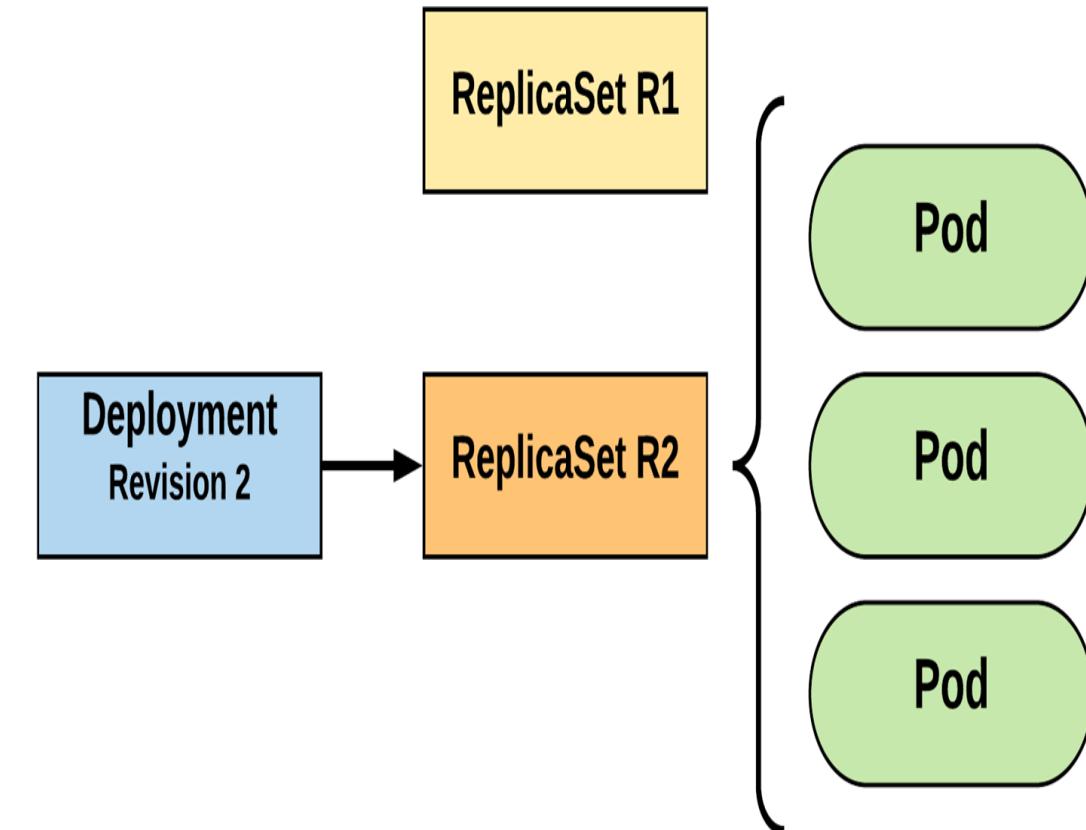
54f7ff7d6d

```
$ kubectl get replicaset
```

NAME	DESIRED	CURRENT	READY	AGE
mydep-54f7ff7d6d	3	3	3	15s
mydep-6766777fff	0	0	0	5h

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
mydep-54f7ff7d6d-9gvll	1/1	Running	0	12s
mydep-54f7ff7d6d-cqvlq	1/1	Running	0	10s
mydep-54f7ff7d6d-gccr6	1/1	Running	0	7s



DaemonSet

- Ensure that all nodes matching certain criteria will run an instance of the supplied Pod.
- They **bypass** default scheduling mechanisms.
- Are ideal for cluster wide services such as log forwarding, or health monitoring.
- Revisions are managed via a controller-revision-hash label.



DaemonSet

- **revisionHistoryLimit**: The number of previous iterations of the DaemonSet to retain.
- **updateStrategy**: Describes the method of updating the Pods based on the **type**. Valid options are **RollingUpdate** or **OnDelete**.
 - **RollingUpdate**: Cycles through updating the Pods according to the value of **maxUnavailable**.
 - **OnDelete**: The new instance of the Pod is deployed **ONLY** after the current instance is deleted.

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: ds-example
spec:
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
  updateStrategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
  template:
    spec:
      nodeSelector:
       .nodeType: edge
<pod template>
```

DaemonSet

- **spec.template.spec.nodeSelector:**
The primary selector used to target nodes.
- **Default Host Labels:**
 - kubernetes.io/hostname
 - beta.kubernetes.io/os
 - beta.kubernetes.io/arch
- **Cloud Host Labels:**
 - failure-domain.beta.kubernetes.io/zone
 - failure-domain.beta.kubernetes.io/region
 - beta.kubernetes.io/instance-type

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: ds-example
spec:
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
  updateStrategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
  template:
    spec:
      nodeSelector:
       .nodeType: edge
      <pod template>
```

DaemonSet

```

apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: ds-example
spec:
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
  updateStrategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
  template:
    metadata:
      labels:
        app: nginx
    spec:
      nodeSelector:
       .nodeType: edge
      containers:
        - name: nginx
          image: nginx:stable-alpine
          ports:
            - containerPort: 80

```

```

$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
ds-example-x8kkz   1/1     Running   0          1m

```

```

$ kubectl describe ds ds-example
Name:         ds-example
Selector:    app=nginx,env=prod
Node-Selector: nodeType=edge
Labels:      app=nginx
             env=prod
Annotations: <none>
Desired Number of Nodes Scheduled: 1
Current Number of Nodes Scheduled: 1
Number of Nodes Scheduled with Up-to-date Pods: 1
Number of Nodes Scheduled with Available Pods: 1
Number of Nodes Misscheduled: 0
Pods Status: 1 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels: app=nginx
          env=prod
  Containers:
    nginx:
      Image:  nginx:stable-alpine
      Port:   80/TCP
      Environment: <none>
      Mounts:  <none>
      Volumes: <none>
  Events:
    Type  Reason          Age   From           Message
    ----  ----          ----  ----          -----
    Normal SuccessfulCreate 48s  daemonset-controller  Created pod: ds-example-x8kkz

```

DaemonSet

```

apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: ds-example
spec:
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
  updateStrategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
  template:
    metadata:
      labels:
        app: nginx
  spec:
    nodeSelector:
     .nodeType: edge
    containers:
      - name: nginx
        image: nginx:stable-alpine
        ports:
          - containerPort: 80

```

NAME	READY	STATUS	RESTARTS	AGE
ds-example-x8kkz	1/1	Running	0	1m

```

$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
ds-example-x8kkz   1/1    Running   0          1m

$ kubectl describe ds ds-example
Name:         ds-example
Selector:    app=nginx,env=prod
Node-Selector: nodeType=edge
Labels:      app=nginx
             env=prod
Annotations: <none>
Desired Number of Nodes Scheduled: 1
Current Number of Nodes Scheduled: 1
Number of Nodes Scheduled with Up-to-date Pods: 1
Number of Nodes Scheduled with Available Pods: 1
Number of Nodes Misscheduled: 0
Pods Status: 1 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=nginx
           env=prod
  Containers:
    nginx:
      Image:  nginx:stable-alpine
      Port:   80/TCP
      Environment: <none>
      Mounts:  <none>
      Volumes: <none>
  Events:
    Type  Reason     Age   From           Message
    ----  -----     --   --    Normal  SuccessfulCreate 48s  daemonset-controller  Created pod: ds-example-x8kkz

```

Job

- Job controller ensures one or more pods are executed and successfully terminate.
- Will continue to try and execute the job until it satisfies the completion and/or parallelism condition.
- Pods are **NOT** cleaned up until the job itself is deleted.*



Job

- **backoffLimit**: The number of failures before the job itself is considered **failed**.
- **completions**: The total number of successful completions desired.
- **parallelism**: How many instances of the pod can be run concurrently.
- **spec.template.spec.restartPolicy**: Jobs only support a **restartPolicy** of type **Never** or **OnFailure**.

```
apiVersion: batch/v1
kind: Job
metadata:
  name: job-example
spec:
  backoffLimit: 4
  completions: 4
  parallelism: 2
  template:
    spec:
      restartPolicy: Never
    <pod-template>
```

```
apiVersion: batch/v1
kind: Job
metadata:
  name: job-example
spec:
  backoffLimit: 4
  completions: 4
  parallelism: 2
  template:
    spec:
      containers:
        - name: hello
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args: ["echo hello from $HOSTNAME!"]
  restartPolicy: Never
```

```
$ kubectl get pods --show-all
NAME      READY   STATUS    RESTARTS   AGE
job-example-dvxd2  0/1   Completed  0          51m
job-example-hknns  0/1   Completed  0          52m
job-example-tphkm  0/1   Completed  0          51m
job-example-v5fvq  0/1   Completed  0          52m
```

```
$ kubectl describe job job-example
Name:      job-example
Namespace: default
Selector:  controller-uid=19d122f4-1576-11e8-a4e2-080027a3682b
Labels:    controller-uid=19d122f4-1576-11e8-a4e2-080027a3682b
           job-name=job-example
Annotations: <none>
Parallelism: 2
Completions: 4
Start Time: Mon, 19 Feb 2018 08:09:21 -0500
Pods Statuses: 0 Running / 4 Succeeded / 0 Failed
Pod Template:
  Labels: controller-uid=19d122f4-1576-11e8-a4e2-080027a3682b
           job-name=job-example
  Containers:
    hello:
      Image: alpine:latest
      Port: <none>
      Command:
        /bin/sh
        -c
      Args:
        echo hello from $HOSTNAME!
      Environment: <none>
      Mounts: <none>
      Volumes: <none>
  Events:
    Type  Reason      Age From      Message
    ----  ----      --  --  -----
    Normal SuccessfulCreate 52m  job-controller  Created pod: job-example-v5fvq
    Normal SuccessfulCreate 52m  job-controller  Created pod: job-example-hknns
    Normal SuccessfulCreate 51m  job-controller  Created pod: job-example-tphkm
    Normal SuccessfulCreate 51m  job-controller  Created pod: job-example-dvxd2
```

- An extension of the Job Controller, it provides a method of executing jobs on a cron-like schedule.

CronJobs within Kubernetes use UTC ONLY.



CronJob

```

apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: cronjob-example
spec:
  schedule: "*/1 * * * *"
  successfulJobsHistoryLimit: 3
  failedJobsHistoryLimit: 1
  jobTemplate:
    spec:
      completions: 4
      parallelism: 2
      template:
        spec:
          containers:
            - name: hello
              image: alpine:latest
              command: ["/bin/sh", "-c"]
              args: ["echo hello from $HOSTNAME!"]
      restartPolicy: Never
  
```

```

$ kubectl get jobs
NAME           DESIRED   SUCCESSFUL   AGE
cronjob-example-1519053240  4        4           2m
cronjob-example-1519053300  4        4           1m
cronjob-example-1519053360  4        4           26s
  
```

```

$ kubectl describe cronjob cronjob-example
Name:           cronjob-example
Namespace:      default
Labels:         <none>
Annotations:   <none>
Schedule:      */1 * * * *
Concurrency Policy: Allow
Suspend:       False
Starting Deadline Seconds: <unset>
Selector:      <unset>
Parallelism:   2
Completions:   4
Pod Template:
  Labels: <none>
  Containers:
    hello:
      Image: alpine:latest
      Port: <none>
      Command:
        /bin/sh
        -c
      Args:
        echo hello from $HOSTNAME!
      Environment: <none>
      Mounts: <none>
      Volumes: <none>
  Last Schedule Time: Mon, 19 Feb 2018 09:54:00 -0500
  Active Jobs:   cronjob-example-1519052040
  Events:
    Type  Reason     Age   From           Message
    ----  -----     --   --            --
    Normal SuccessfulCreate 3m  cronjob-controller  Created job cronjob-example-1519051860
    Normal SawCompletedJob 2m  cronjob-controller  Saw completed job: cronjob-example-1519051860
    Normal SuccessfulCreate 2m  cronjob-controller  Created job cronjob-example-1519051920
    Normal SawCompletedJob 1m  cronjob-controller  Saw completed job: cronjob-example-1519051920
    Normal SuccessfulCreate 1m  cronjob-controller  Created job cronjob-example-1519051980
  
```

Storage - Concepts and Resources

- Storage
- Volumes
- Persistent Volumes
- Persistent Volume Claims

- Pods by themselves are useful, but many workloads require exchanging data between containers, or persisting some form of data.
- For this we have **Volumes**, **PersistentVolumes**, **PersistentVolumeClaims**, and **StorageClasses**.

Volumes

- Storage that is tied to the **Pod's Lifecycle**.
- A pod can have one or more types of volumes attached to it.
- Can be consumed by any of the containers within the pod.
- Survive Pod restarts; however their durability beyond that is dependent on the Volume Type.

Volume Types

- awsElasticBlockStore
- azureDisk
- azureFile
- cephfs
- configMap
- csi
- downwardAPI
- emptyDir
- fc (fibre channel)
- flocker
- gcePersistentDisk
- gitRepo
- glusterfs
- hostPath
- iscsi
- local
- nfs
- persistentVolumeClaim
- projected
- portworxVolume
- quobyte
- rbd
- scaleIO
- secret
- storageos
- vsphereVolume



Persistent Volume Supported

Volumes

- **volumes**: A list of volume objects to be attached to the Pod. Every object within the list must have it's own unique **name**.
- **volumeMounts**: A container specific list referencing the Pod volumes by **name**, along with their desired **mountPath**.

```
apiVersion: v1
kind: Pod
metadata:
  name: volume-example
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
      volumeMounts:
        - name: html
          mountPath: /usr/share/nginx/html
          ReadOnly: true
    - name: content
      image: alpine:latest
      command: ["/bin/sh", "-c"]
      args:
        - while true; do
            date >> /html/index.html;
            sleep 5;
        done
      volumeMounts:
        - name: html
          mountPath: /html
  volumes:
    - name: html
      emptyDir: {}
```

- **volumes**: A list of volume objects to be attached to the Pod. Every object within the list must have it's own unique **name**.
- **volumeMounts**: A container specific list referencing the Pod volumes by **name**, along with their desired **mountPath**.

```
apiVersion: v1
kind: Pod
metadata:
  name: volume-example
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
      volumeMounts:
        - name: html
          mountPath: /usr/share/nginx/html
          ReadOnly: true
        - name: content
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args:
            - while true; do
                date >> /html/index.html;
                sleep 5;
            done
      volumeMounts:
        - name: html
          mountPath: /html
  volumes:
    - name: html
      emptyDir: {}
```

- **volumes**: A list of volume objects to be attached to the Pod. Every object within the list must have it's own unique **name**.
- **volumeMounts**: A container specific list referencing the Pod volumes by **name**, along with their desired **mountPath**.

```
apiVersion: v1
kind: Pod
metadata:
  name: volume-example
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
      volumeMounts:
        - name: html
          mountPath: /usr/share/nginx/html
          readOnly: true
    - name: content
      image: alpine:latest
      command: ["/bin/sh", "-c"]
      args:
        - while true; do
            date >> /html/index.html;
            sleep 5;
        done
      volumeMounts:
        - name: html
          mountPath: /html
  volumes:
    - name: html
      emptyDir: {}
```

THANK YOU

